


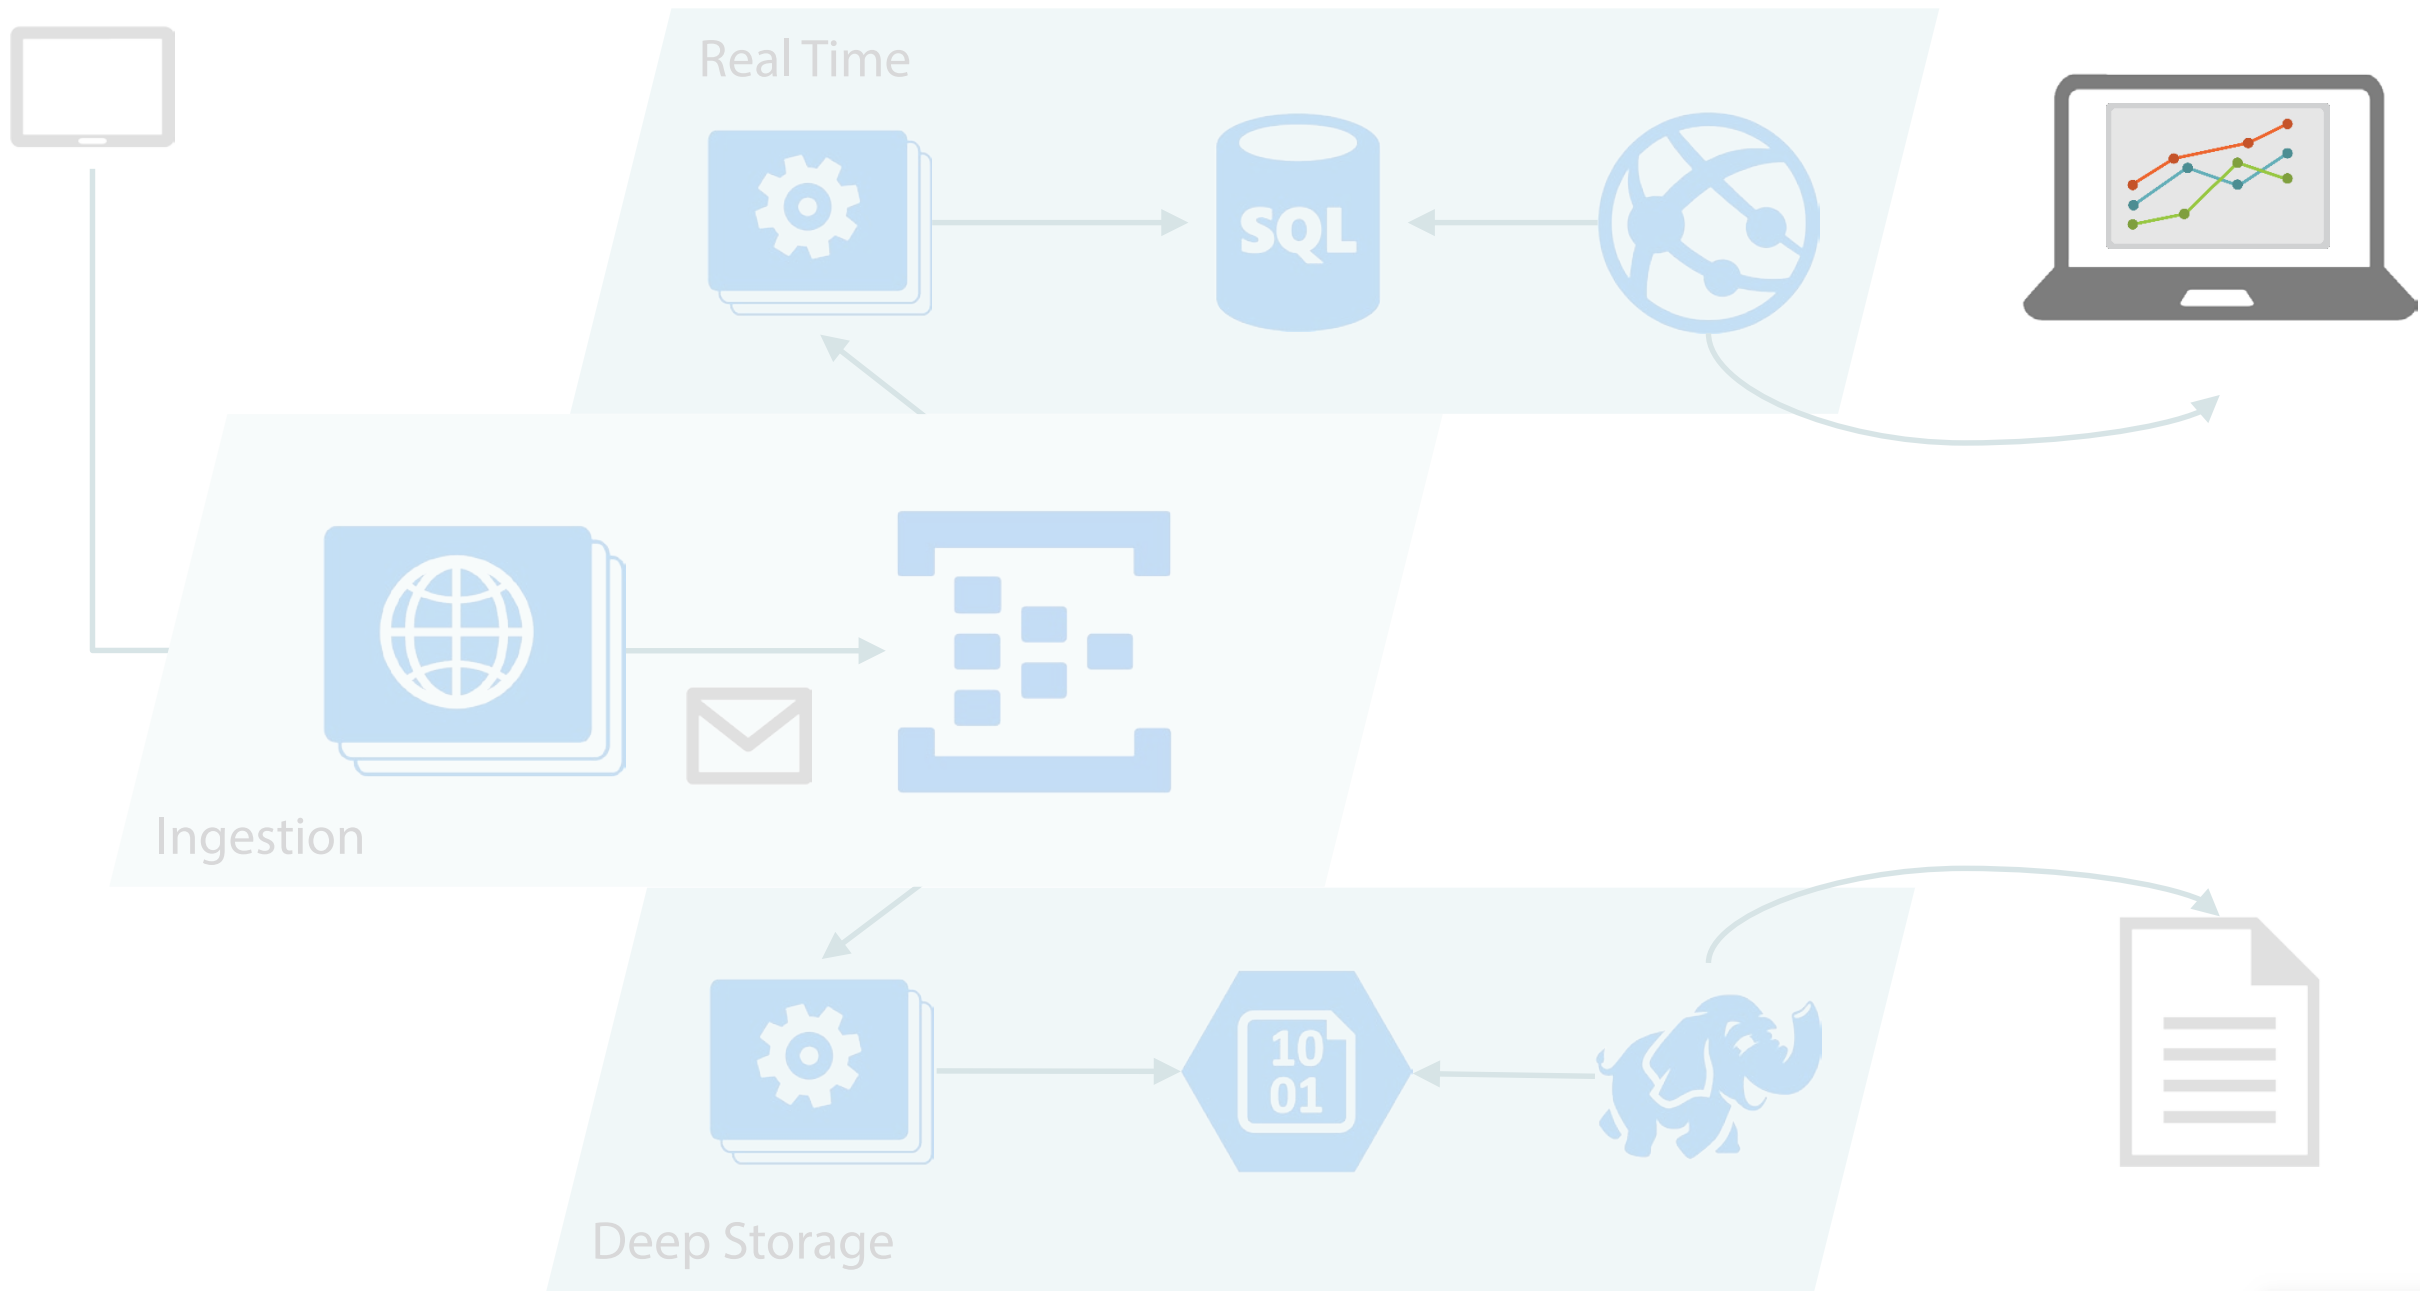
# Storing Event Data for Batch Queries



Elton Stoneman

@EltonStoneman | [www.geekswithblogs.net/eltonstoneman](http://www.geekswithblogs.net/eltonstoneman)





Real Time

Current Firmware  
**V1.2.0RC5**

rel.android-build.20150106.154003  
Last updated at 22:08:27

User Base

**57.4%**

316222 out of 554288 devices  
Last updated at 22:07:58



Last updated at 22:08:01

Latest BIOS  
**44.0**

BayTrail.5.04.16.0044.0  
Last updated at 22:08:27

User Base

**0%**

1 out of 554289 devices  
Last updated at 22:08:28



Last updated at 22:08:00



Ing


Deep Storage

Real Time

Today's Events


506.8M

Last updated at 22:09.27



gained  
lost  
updated  
installed  
removed

Last updated at 22:09.28




switched  
updated  
created  
deleted

Last updated at 22:09.28

Top Events


device.logs.message	409.8M
device.metrics.sampled	28.0M
system.app.focus.gained	12.6M
system.app.focus.lost	12.2M
device.gps.usage.stopped	5.9M
(other)	38.4M

Last updated at 22:09.28



logs.message  
metrics.sampled  
app.focus.gained  
app.focus.lost  
gps.usage.stopped  
(other)

Last updated at 22:09.28

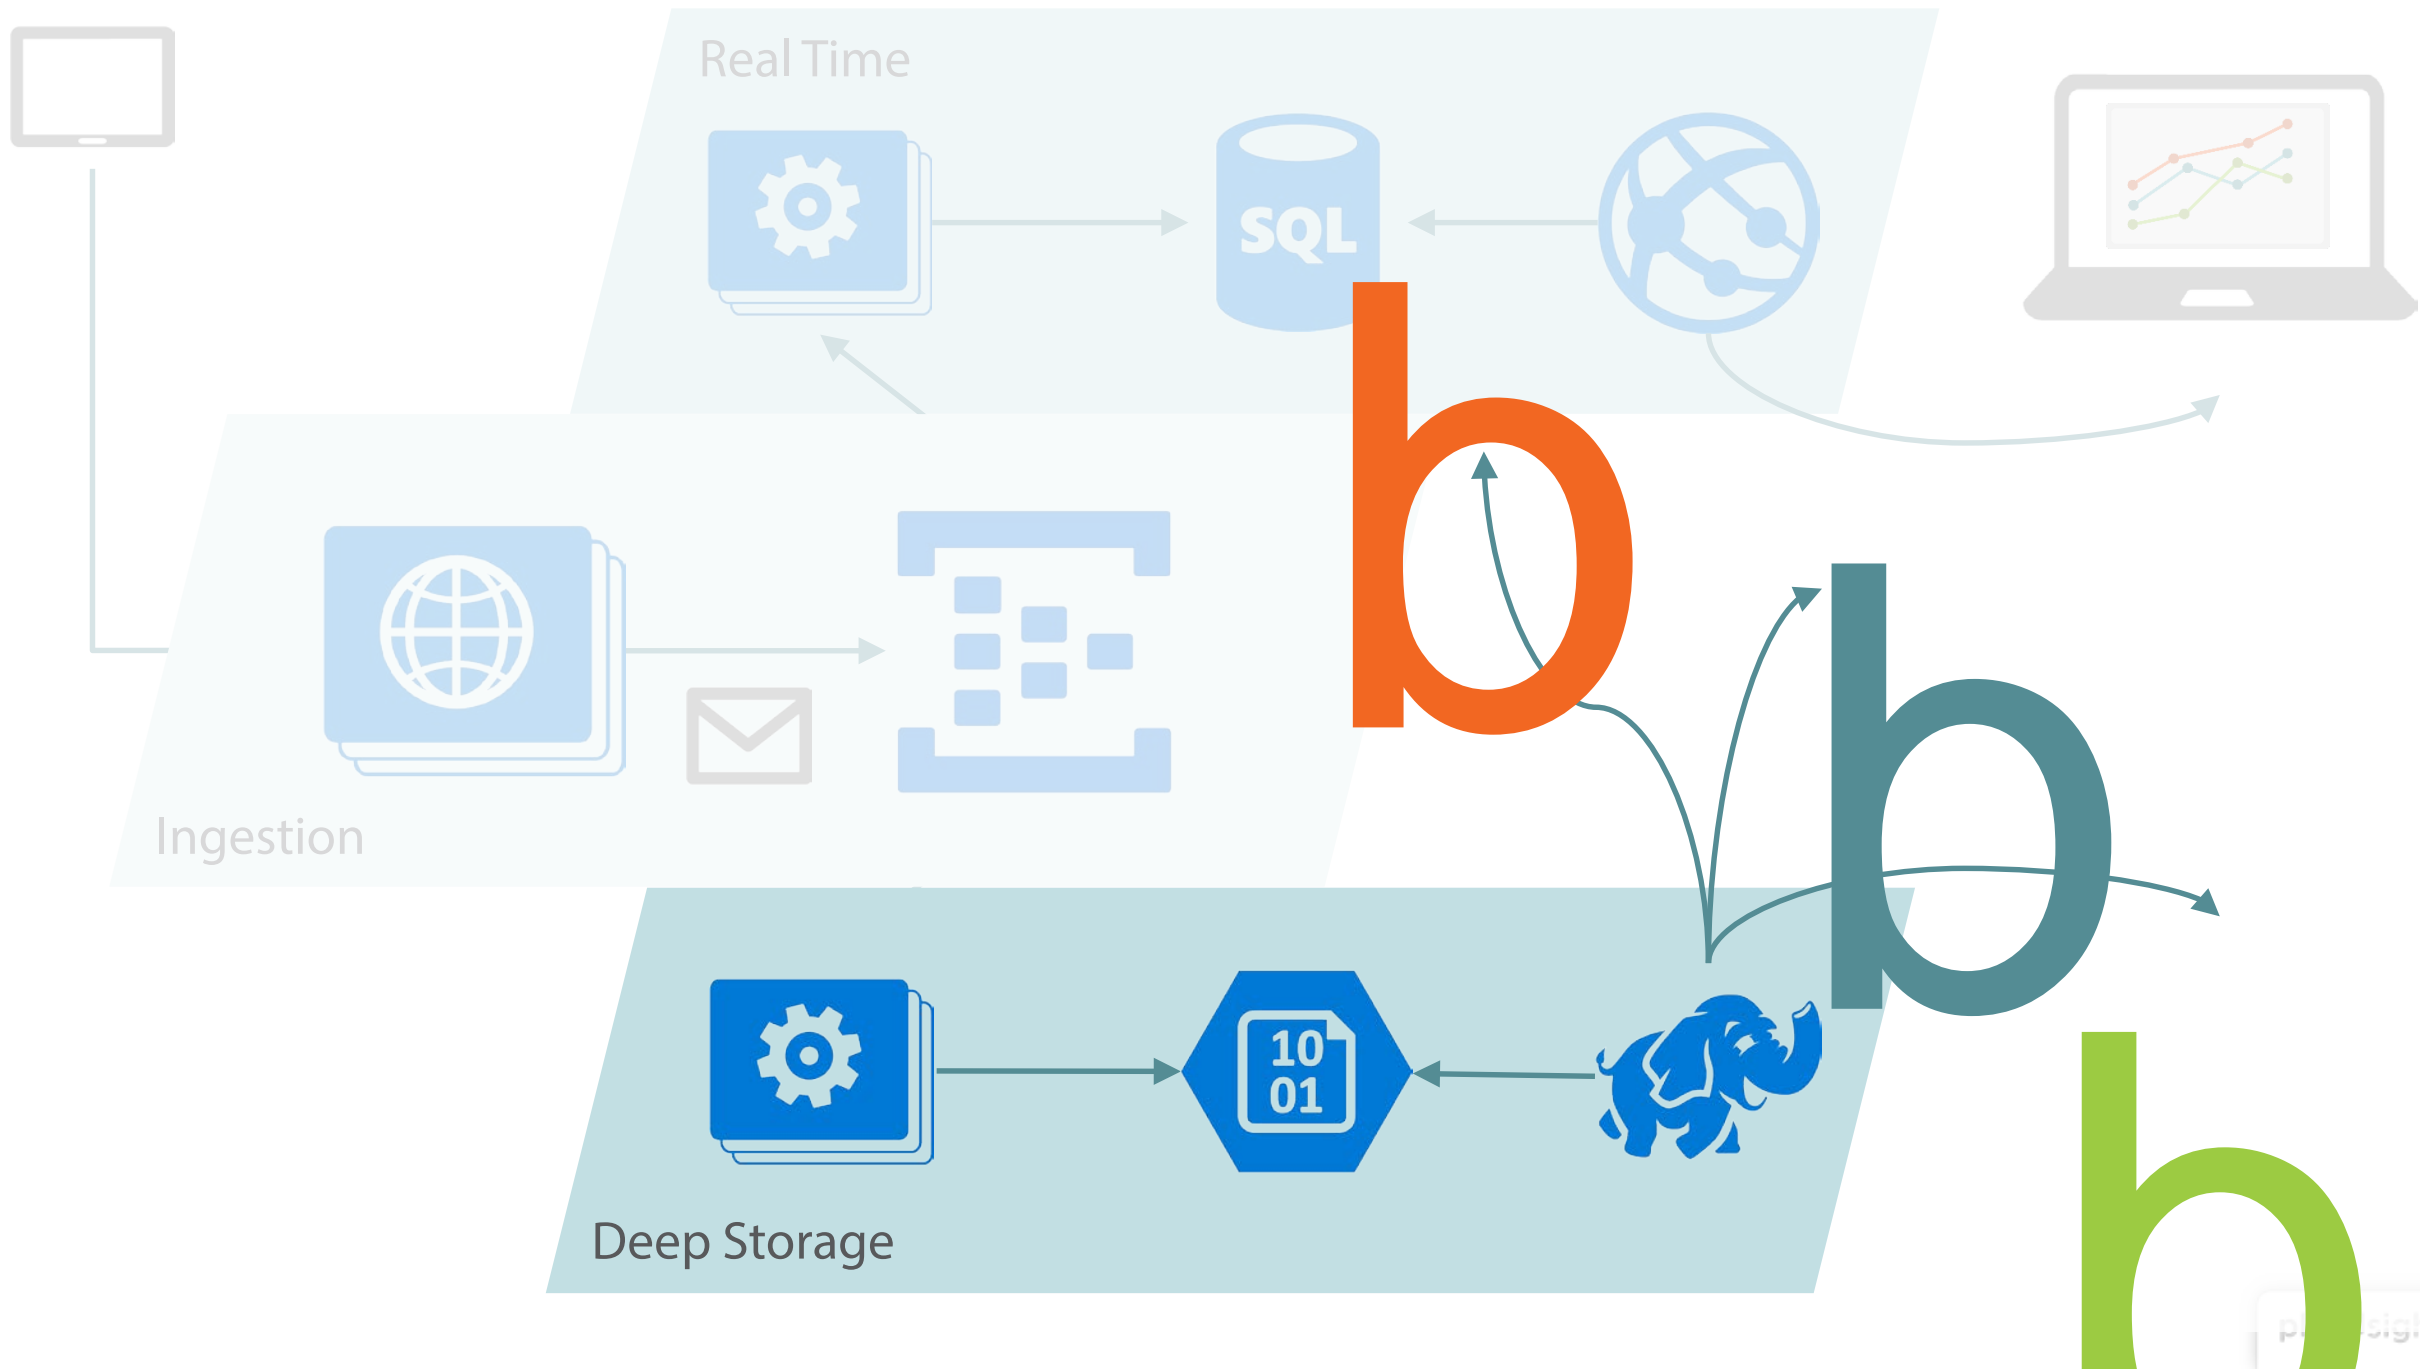


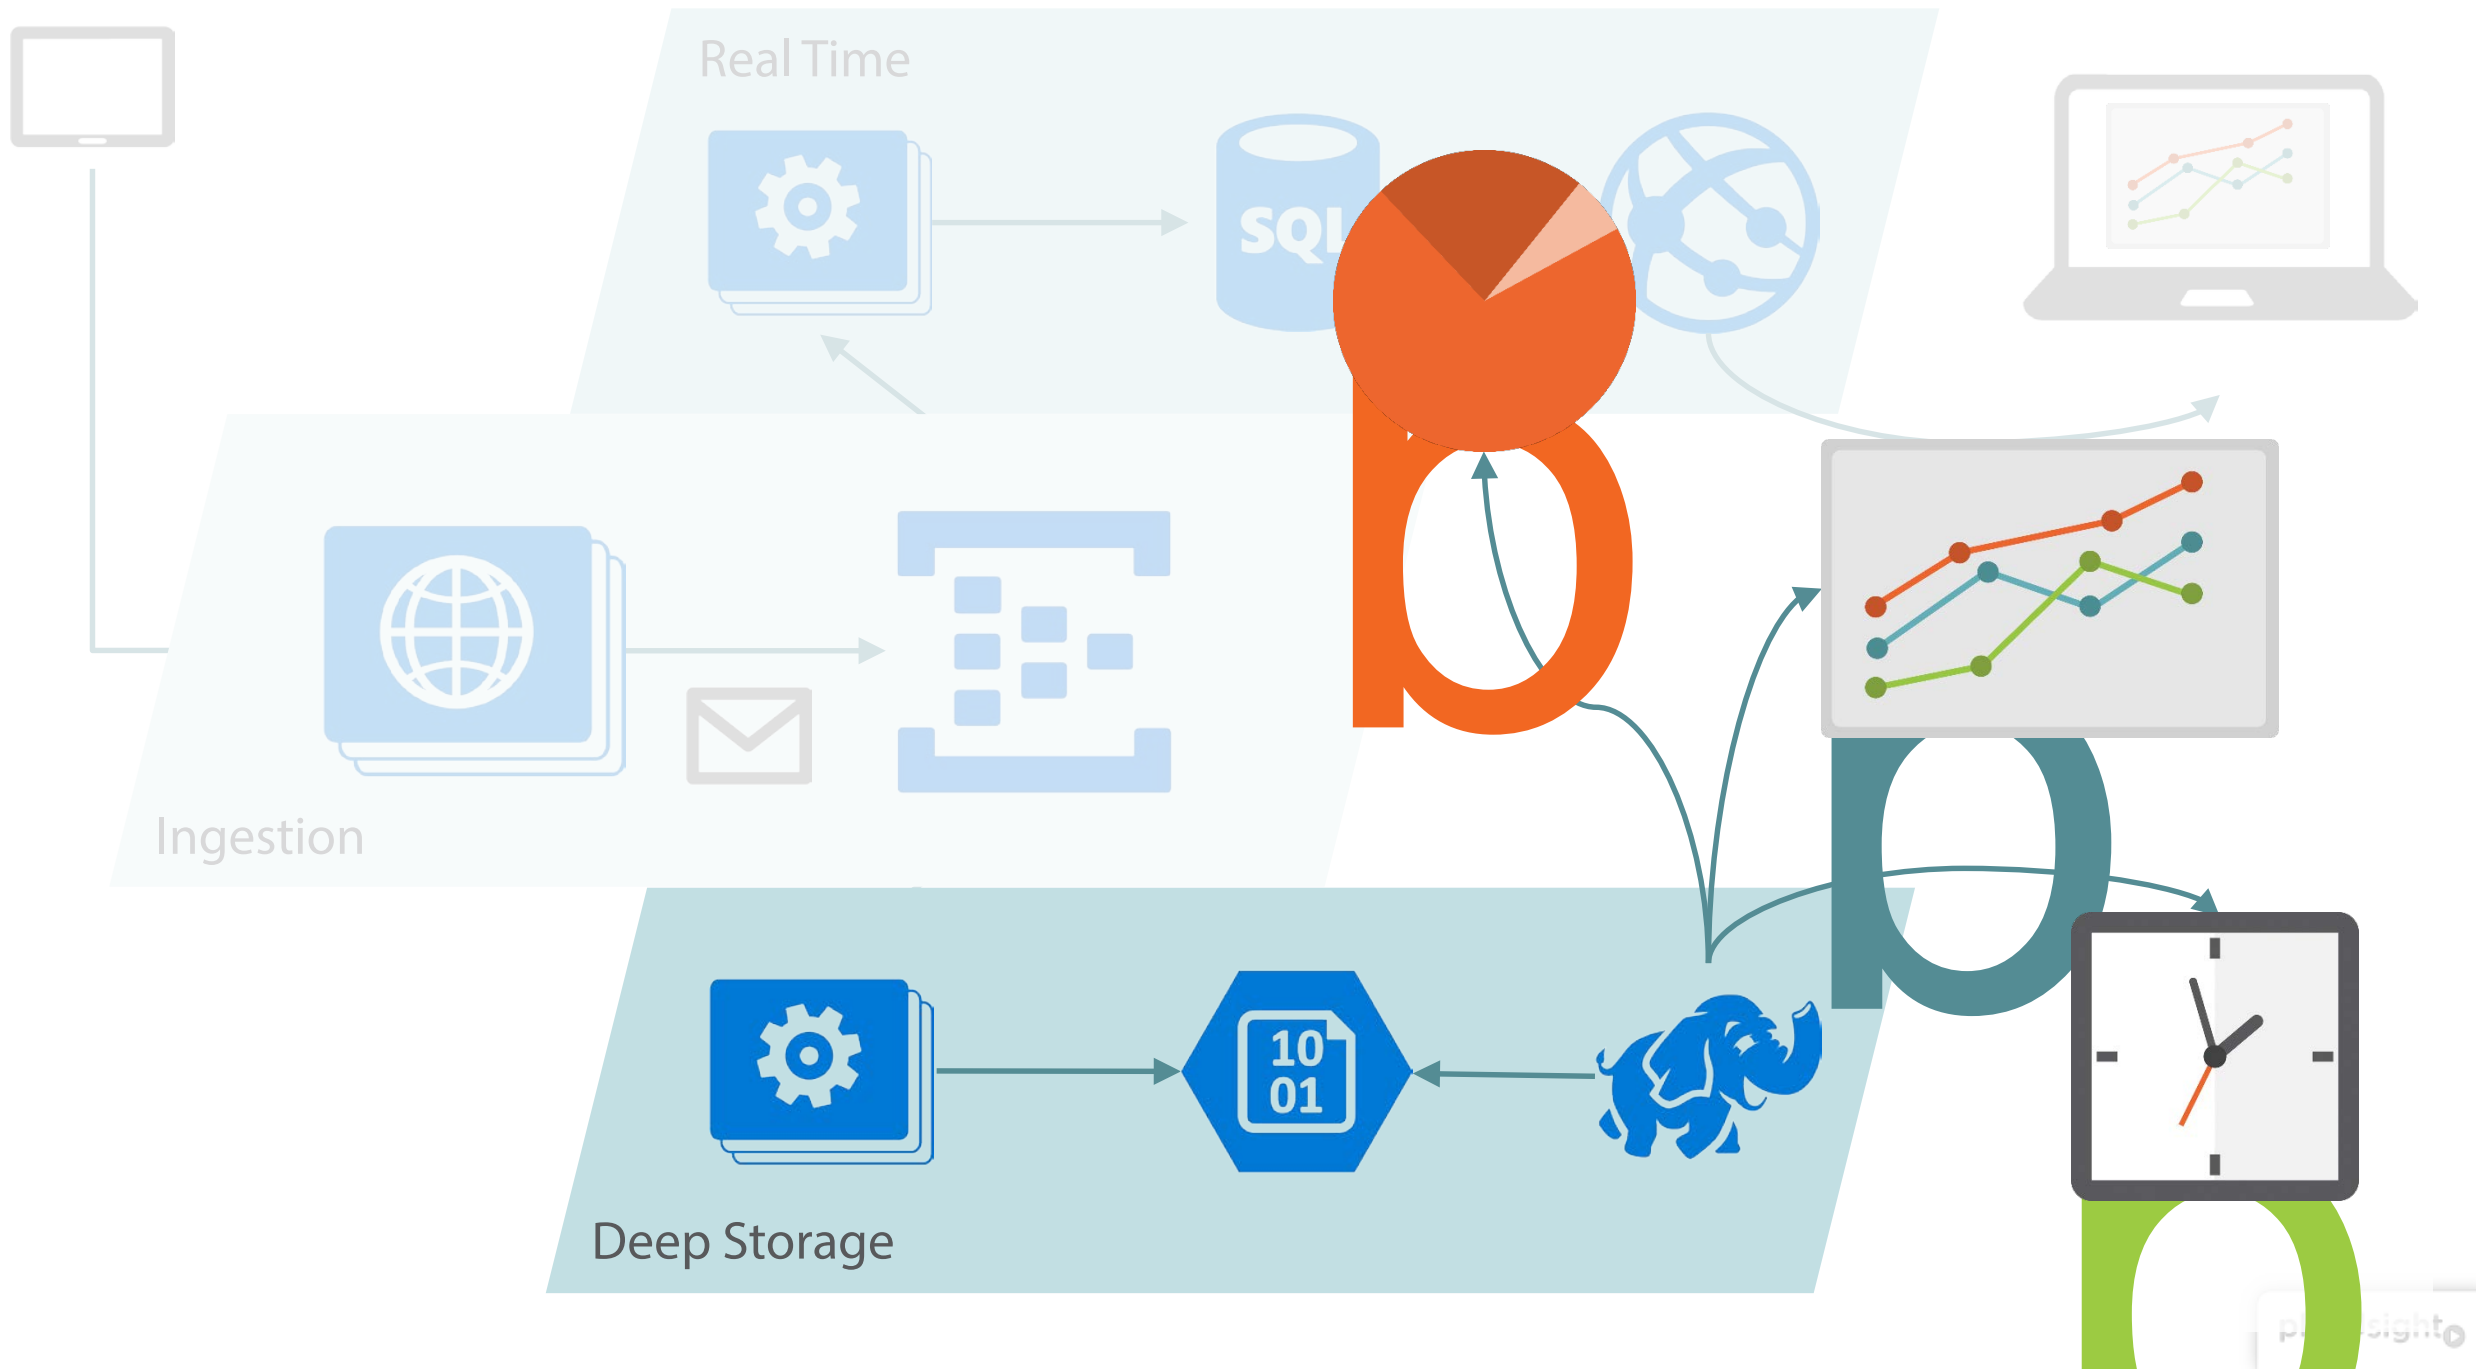
connected  
disconnected  
enabled  
disabled

Last updated at 22:09.28

Deep Storage









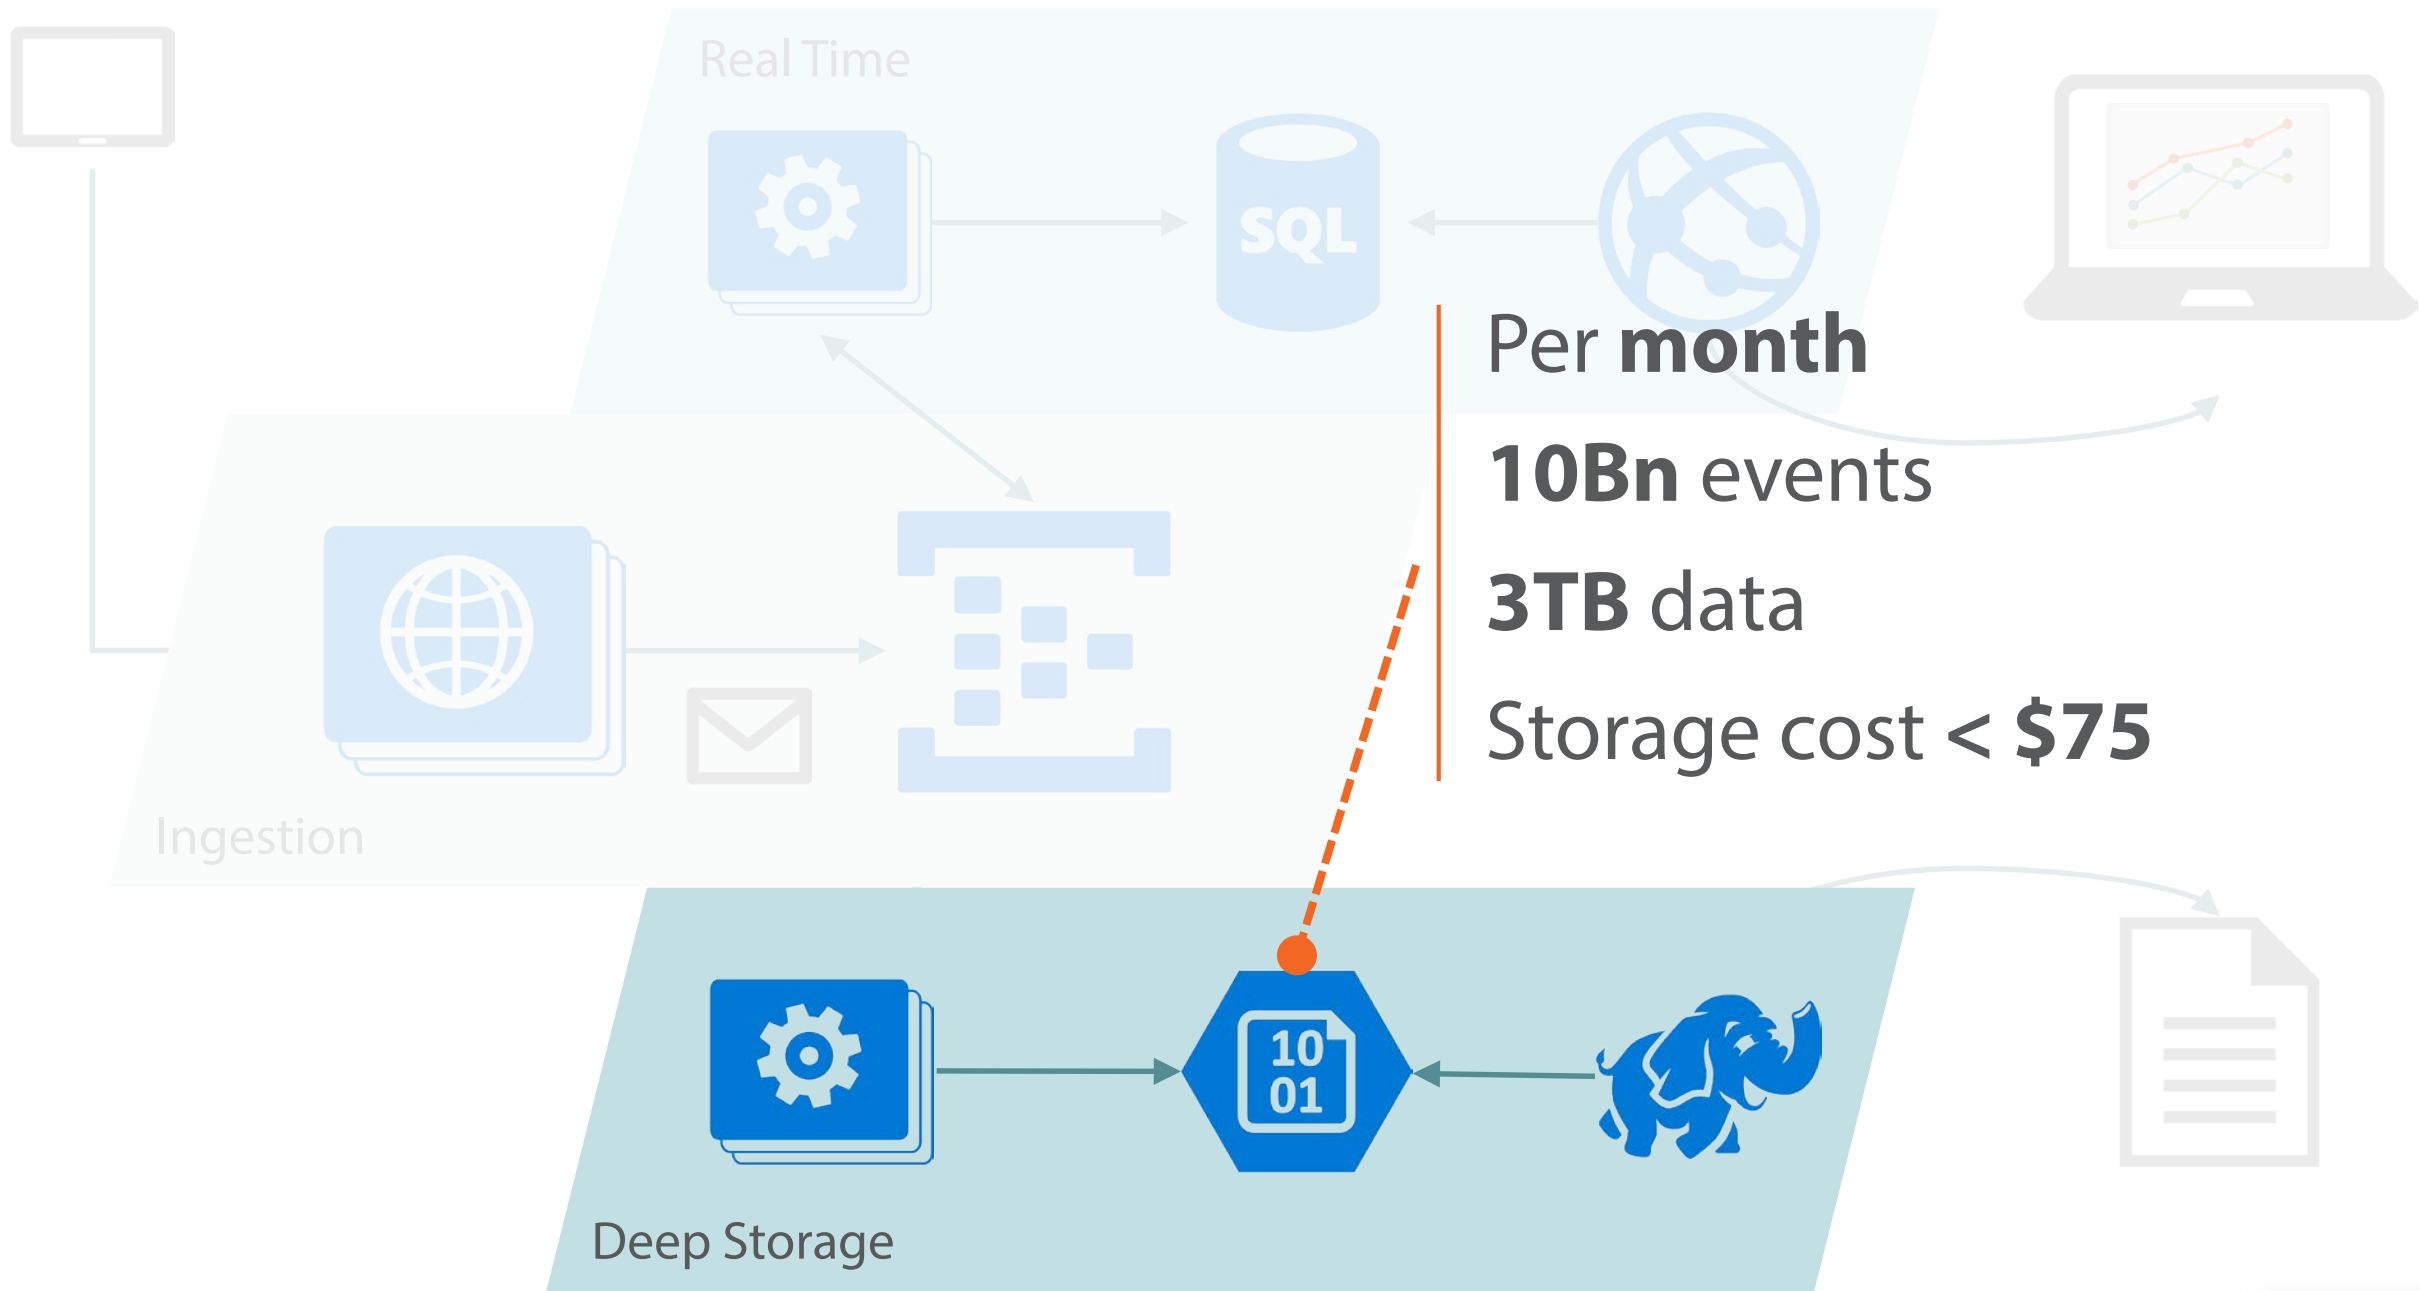
## Azure Blob Storage

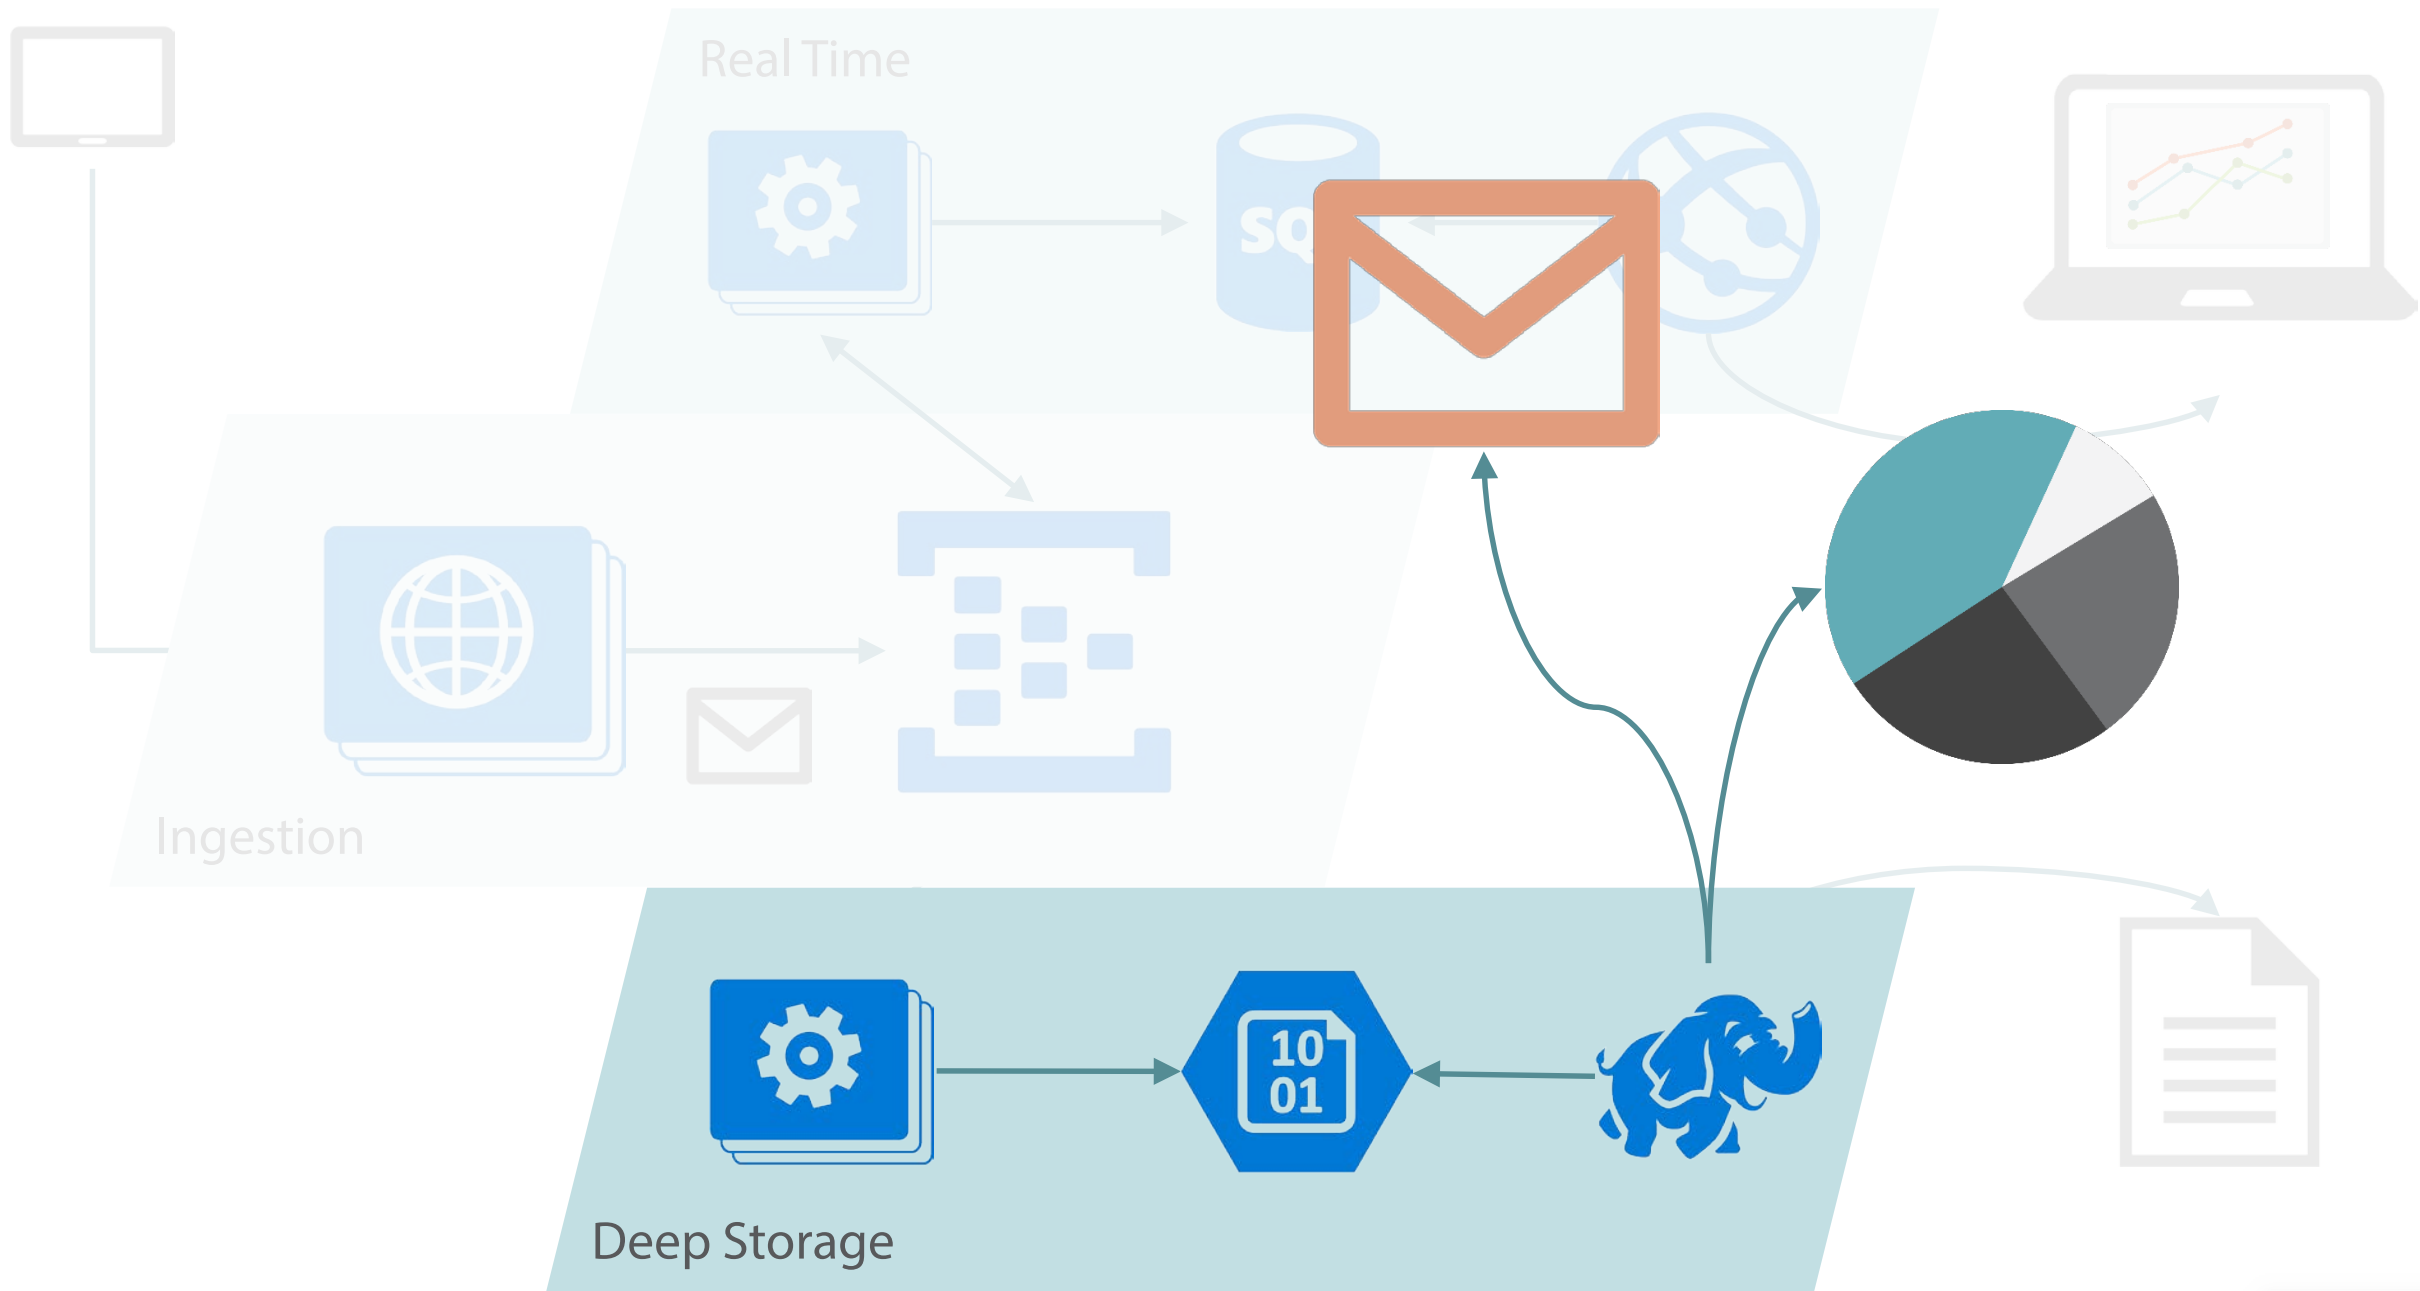
Hadoop HDFS Compatible


Local & Geo-Replication

Billed at < \$25 per TB












- Efficiency
- Reliability
- Traceability



{message-id}.json



Raw **JSON**  
<1KB per blob  
**500M** per day



\*.json




...





x BILLIONs




\*.json.gz





x THOUSANDs




`{partition}/{hour}.json.gz`








`{partition}/{hour}.json.gz`









`{partition}/{hour}.json.gz`





`{partition}/{hour}.json.gz`




**GZip** compressed  
**20-100MB** per blob  
**384** per day



\*.json.gz





...



**384 files**

**300GB JSON**


**500M events**



`{partition}/{hour}.json.gz`



**GZip** compressed  
**20-100MB** per blob  
**384** per day




{hour}.json.gz





**GZip** compressed  
**0.3-1.6GB** per blob  
**24** per day

Peak **40M** per hour  
>**600K** per minute  
>**11K** per second




Peak **40M** per hour  
>**600K** per minute  
>**11K** per second












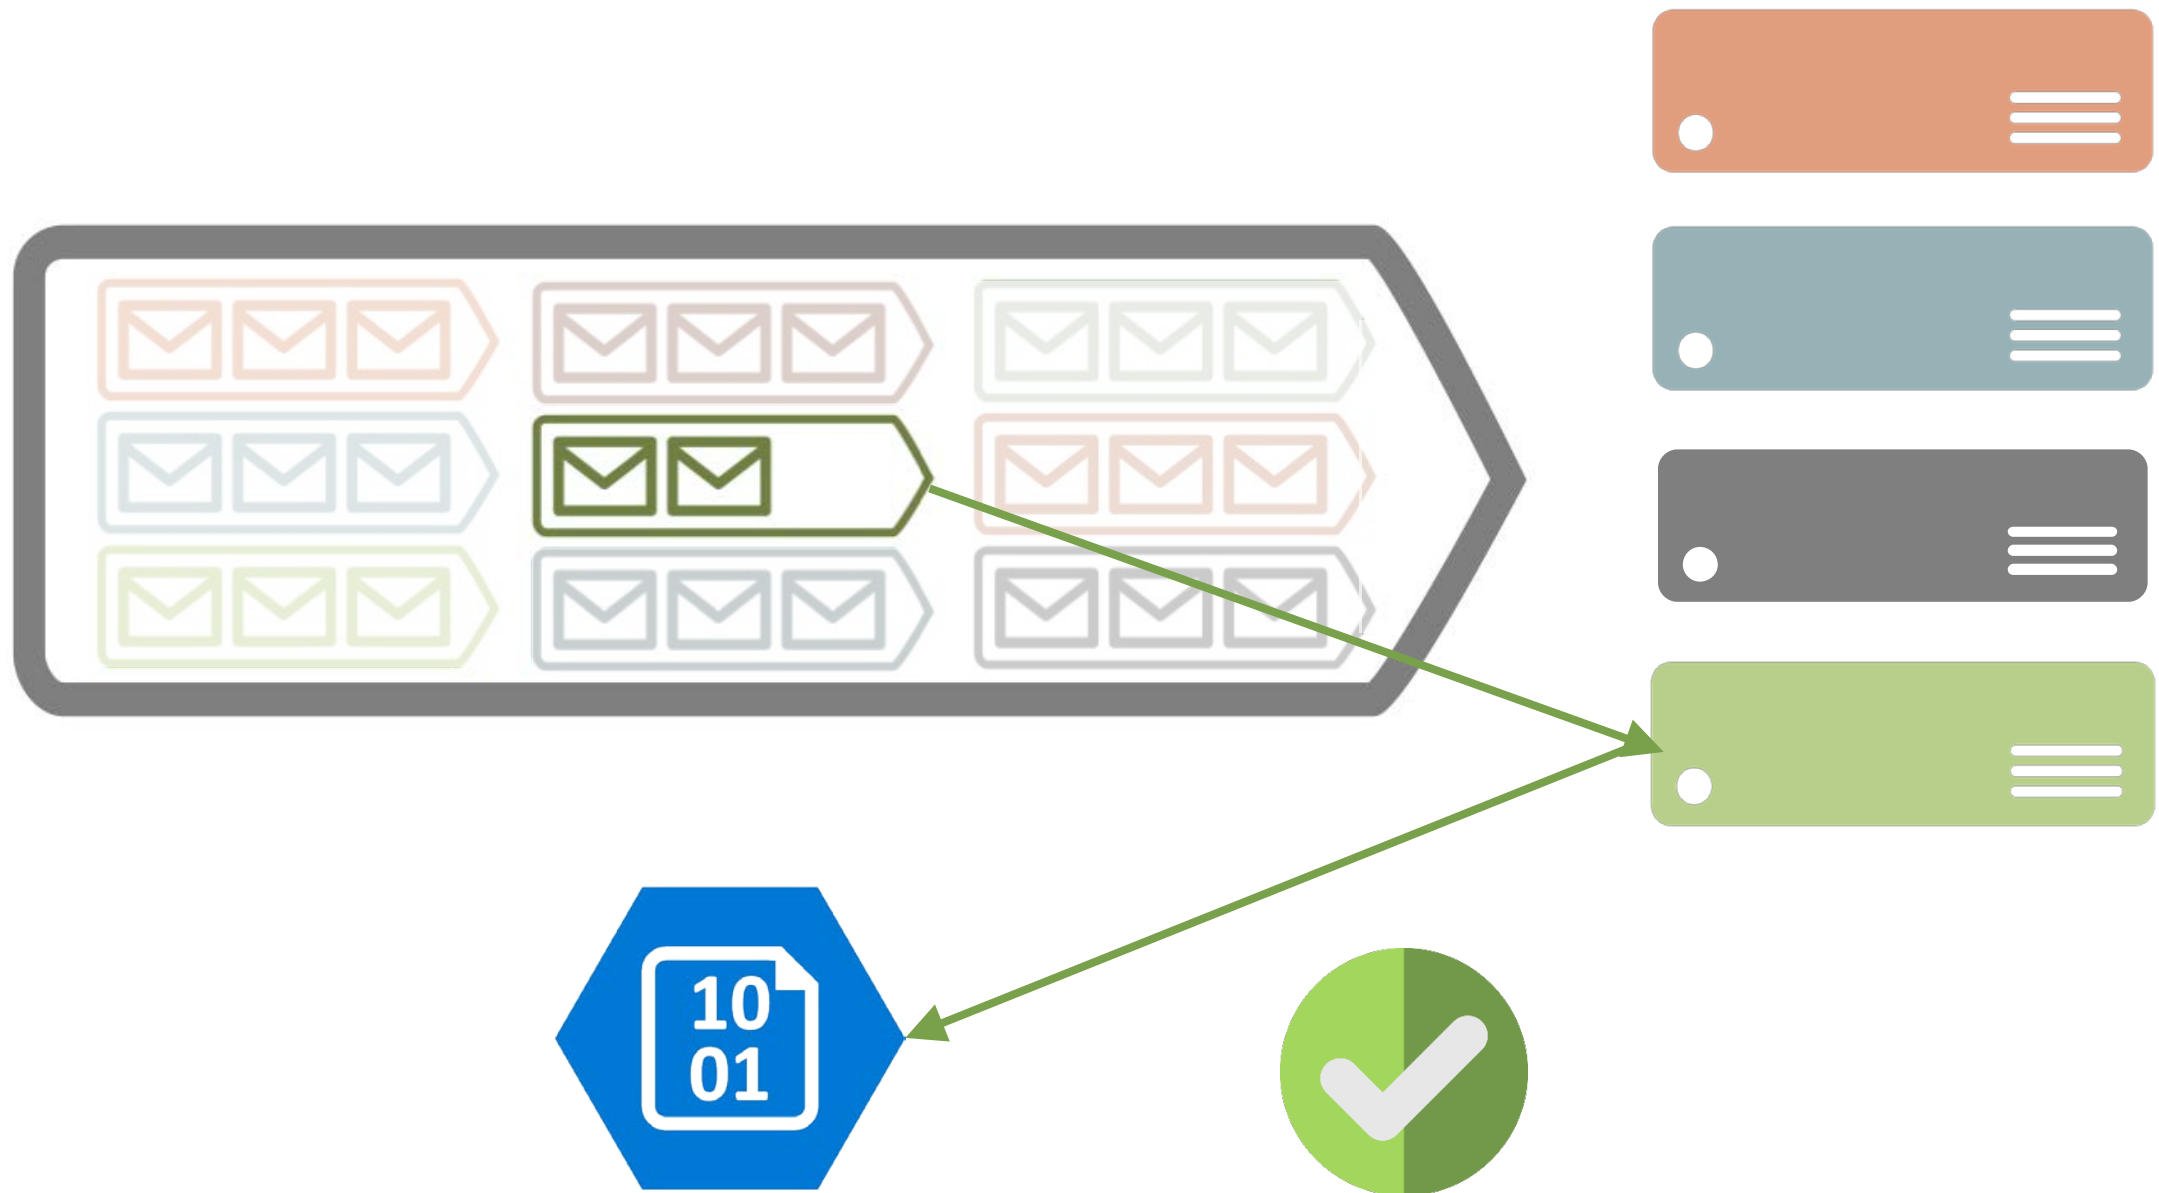
**{partition}/{hour}.json.gz**

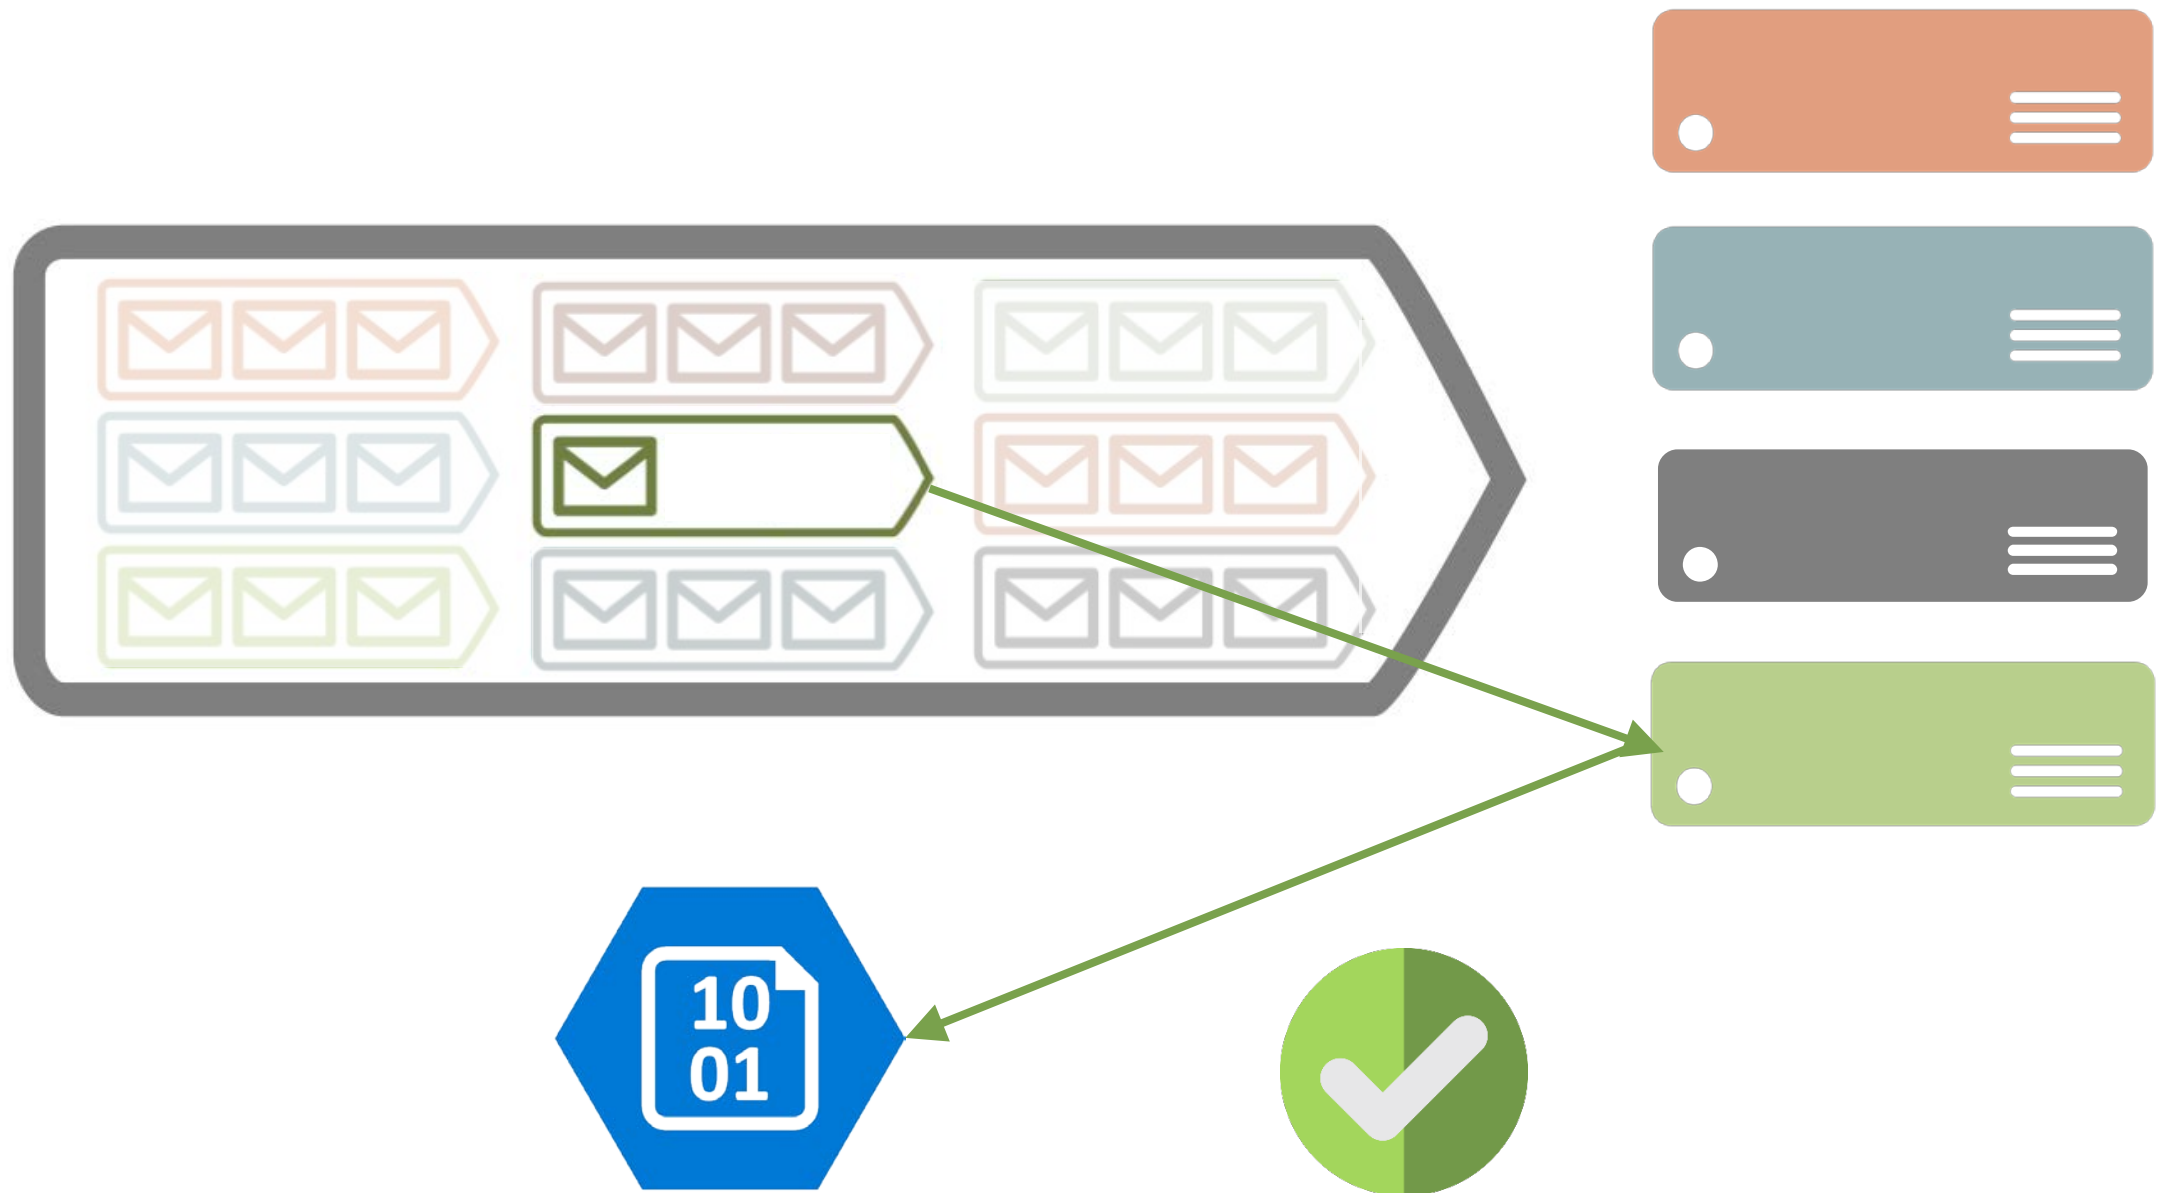


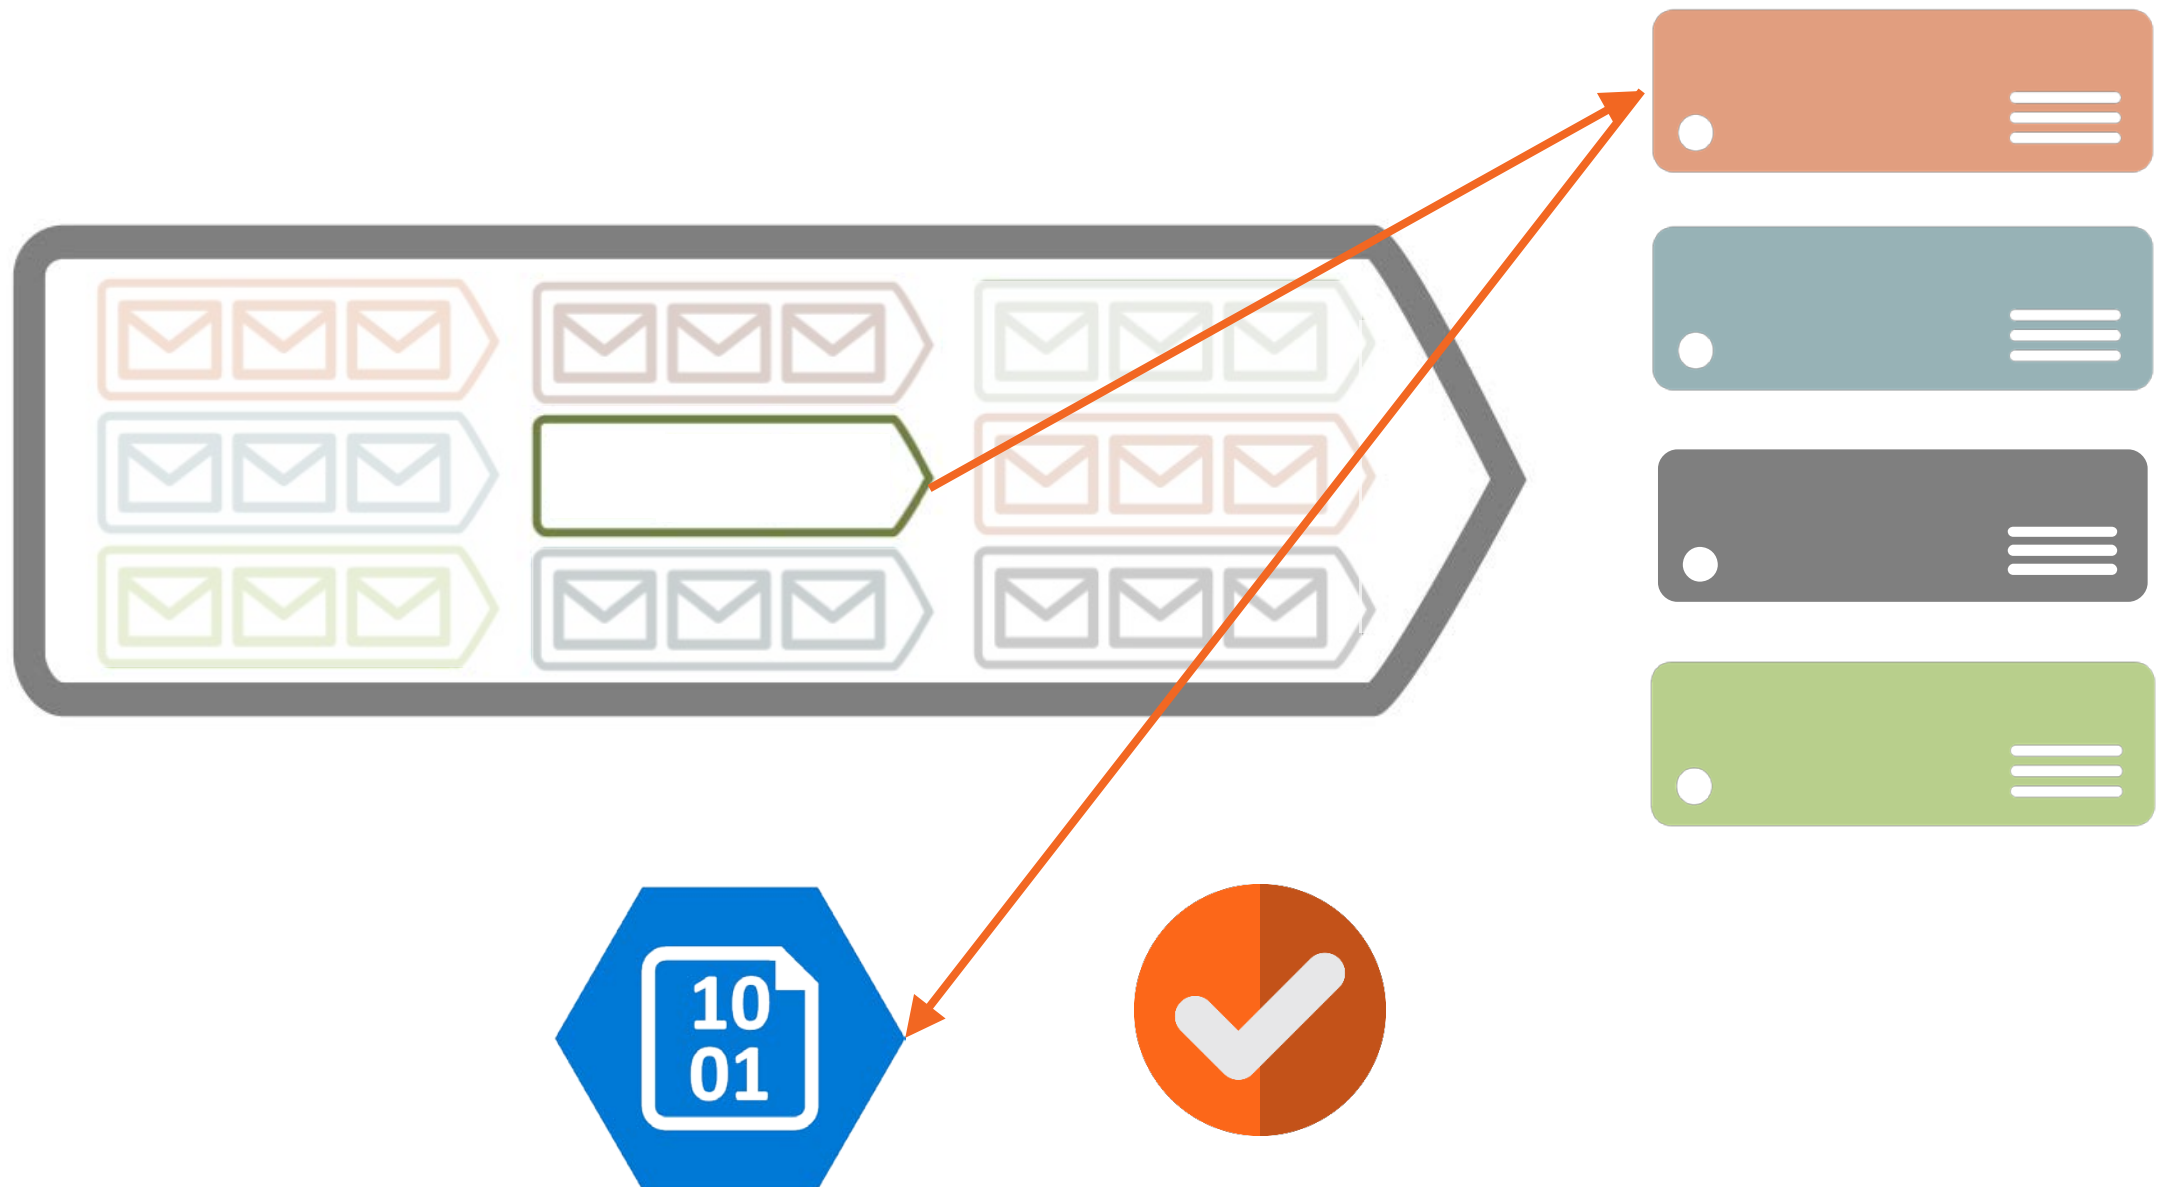




**EventProcessorHost**



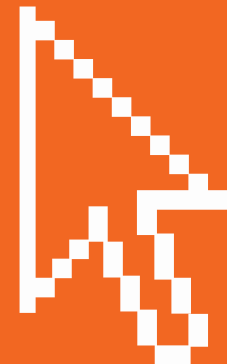




# Demo: IEventProcessor

EventReceiver gets events

IEventProcessor handles events





```
_eventProcessorHost = new EventProcessorHost(  
    _hostName, _eventHubName, _consumerGroupName,  
    _hubConnectionString, _checkpointConnectionString);  
  
await _eventProcessorHost.RegisterEventProcessorAsync  
    <DeepStorageEventProcessor>(processorOptions);
```

## EventReceiver

Initiaize EventProcessorHost & register IEventProcessor

```
var processorOptions = new EventProcessorOptions
{
    MaxBatchSize = 5000,
    PrefetchCount = 1000
};
```

## EventProcessorOptions


Specify receiver throughput

```
Task OpenAsync(PartitionContext context);  
Task CloseAsync(PartitionContext context,  
                CloseReason reason);  
Task ProcessEventsAsync(PartitionContext context,  
                        IEnumerable<EventData> messages);
```


## IEventProcessor


Stateful processor for event batches





```
{  
  "PartitionId": "10",  
  "Owner": "RD000D3AB06D27",  
  "Token": "a2d958f7-e909...",  
  "Epoch": 2,  
  "Offset": "",  
  "SequenceNumber": 0  
}
```

















`{partition}/{hour}.json.gz`






`{partition}/{hour}.json.gz`






`{partition}/{hour}.json.gz`





`{partition}/{hour}.json.gz`




`{partition}/{hour}.json.gz`




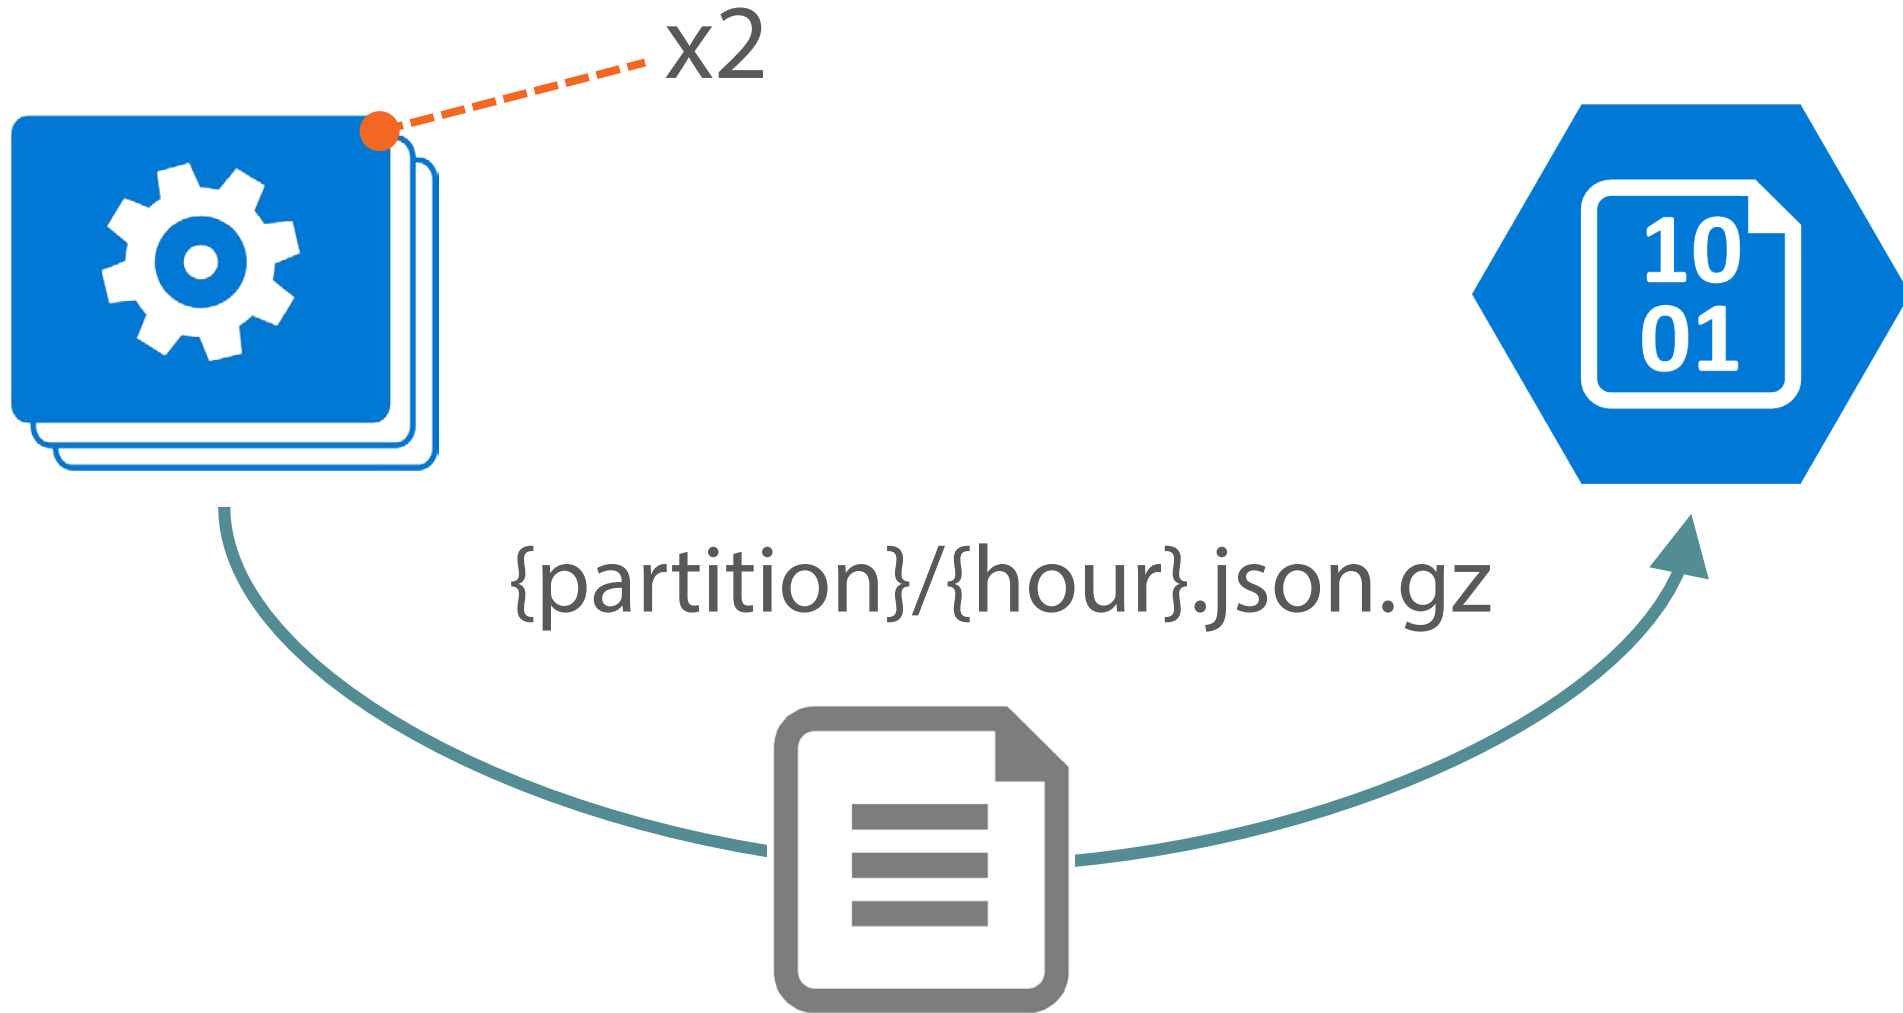



`{partition}/{hour}.json.gz`


















2x **small** instances  
**500M** per day  
Avg **50%** CPU




`{partition}/{hour}.json.gz`





16x **small** instances  
4Bn overnight  
90+% CPU



{partition}/{hour}.json.gz

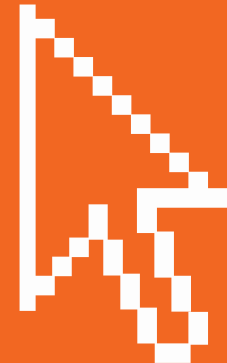


# Demo: DeepStorageEventProcessor

IEventStore interface

Multiple buffer abstraction

Event Processor -> IEventStore



```
foreach (EventData eventData in messages)
{
    var store = GetEventStore(eventData, partitionId);
    var bytes = eventData.GetBytes();
    store.Write(bytes);
}
```

## DeepStorageEventProcessor

Write events to Level 1 Event Store




```
var key = string.Format("{0}p{1}", receivedAt, partitionId);  
if (!_EventStores.ContainsKey(key)) {  
    var store = Container.Instance.Resolve<IEventStore>("1");  
    store.Initialise(partitionId, receivedAt);  
    _EventStores[key] = store;  
}
```

## DeepStorageEventProcessor

Event Stores in static ConcurrentDictionary

DeepStorage

MemoryEventStore  
Write(byte[])



L1

DeepStorage


MemoryEventStore

Write(byte[])

Flush()

DiskEventStore

Write(byte[])



L1 8MB



L2

DeepStorage


MemoryEventStore

Write(byte[])

Flush()

DiskEventStore

Write(byte[])



**L1 8MB**

**16 partitions**

**2 hours**

**256MB L1 buffer**


L2

DeepStorage

MemoryEventStore

Write(byte[])

? \_buffer



L1

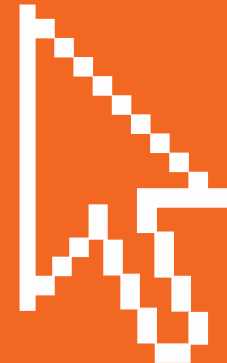
~~Stream~~  
~~byte[]~~  
StringBuilder

# Demo: MemoryEventStore

IEventStore implementation

In-memory store

StringBuilder buffer



```
var json = Encoding.UTF8.GetString(value);  
try  
{  
    _lock.Wait();  
    _buffer.AppendLine(json);  
}
```

## MemoryEventStore

Store events in StringBuilder

```
if (_buffer.Length + byteCount > MaxBufferSize)
{
    Flush();
}
```

## MemoryEventStore

Flush before buffer size exceeded



```
try
{
    _lock.Wait();
    block = _buffer.ToString();
    _buffer.Clear();
}
```

## MemoryEventStore

Read whole StringBuilder on flush

```
if (block.Length > 0)
{
    var data = Encoding.UTF8.GetBytes(block);
    Task.Factory.StartNew(() => _nextStore.Write(data));
}
```

## MemoryEventStore

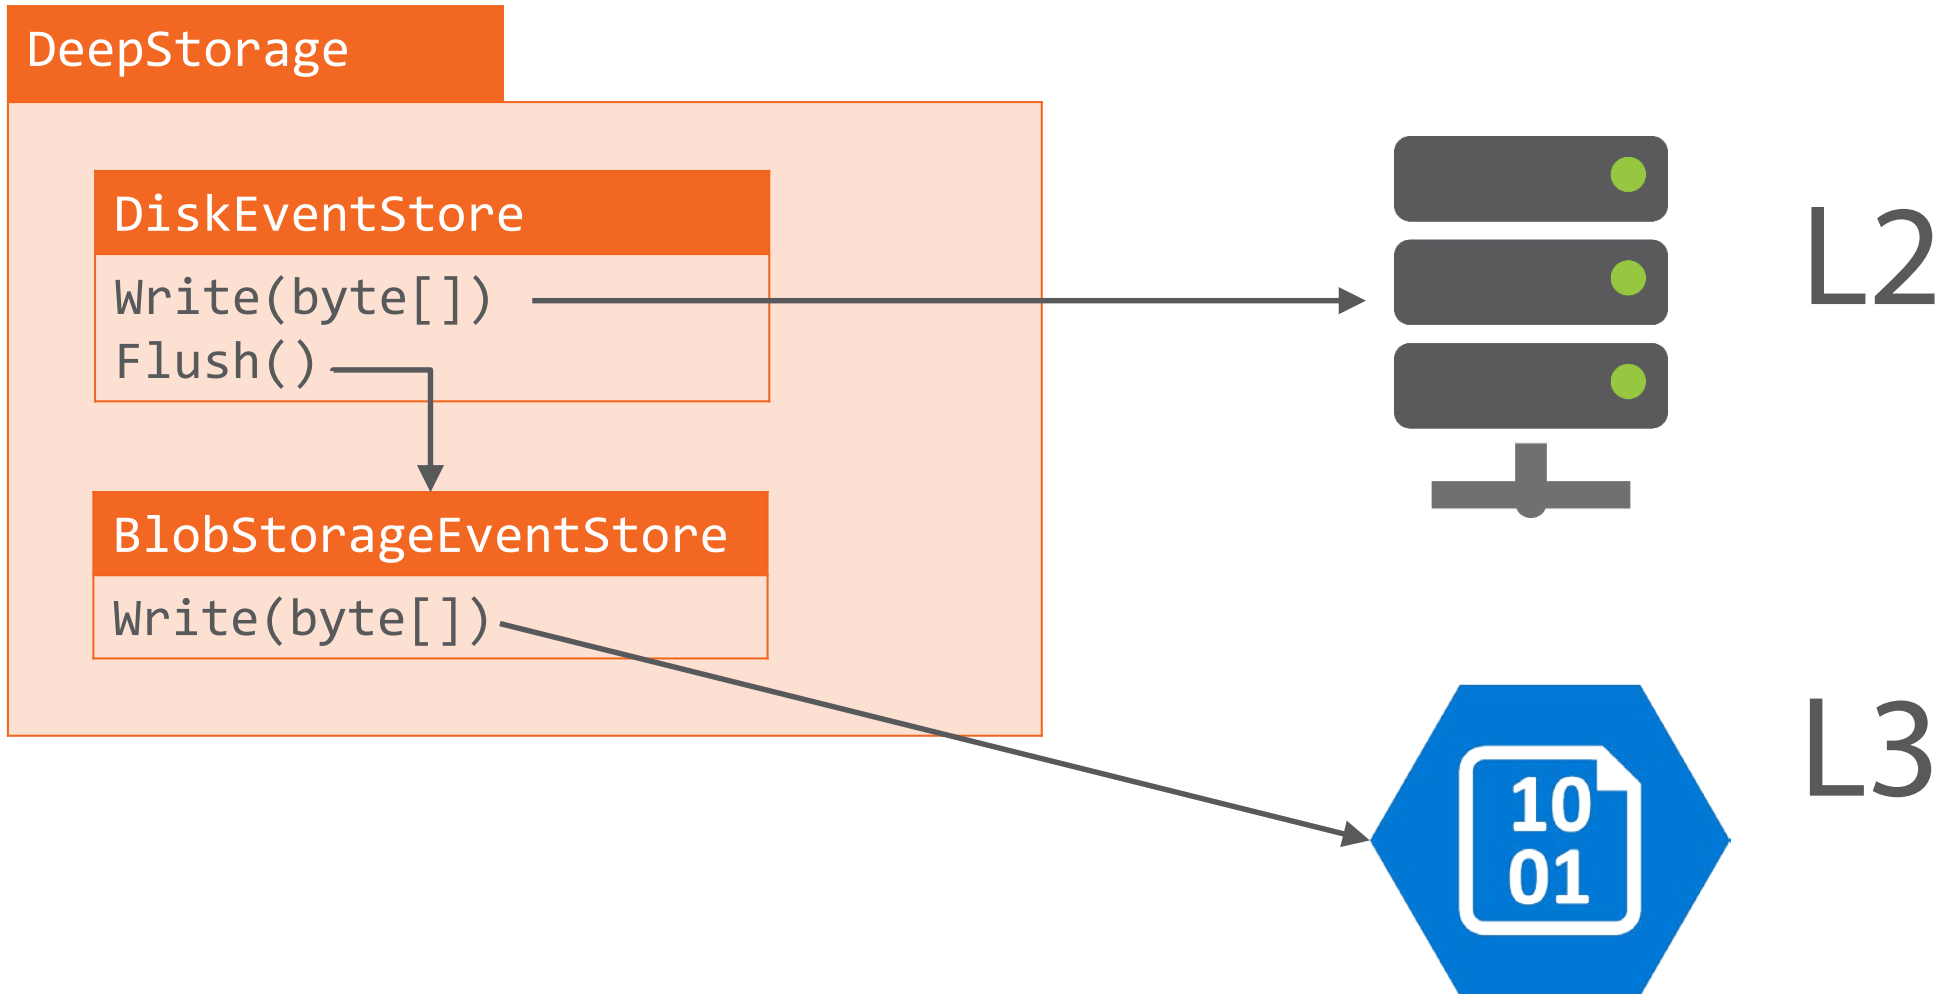
New task, writing to L2 store

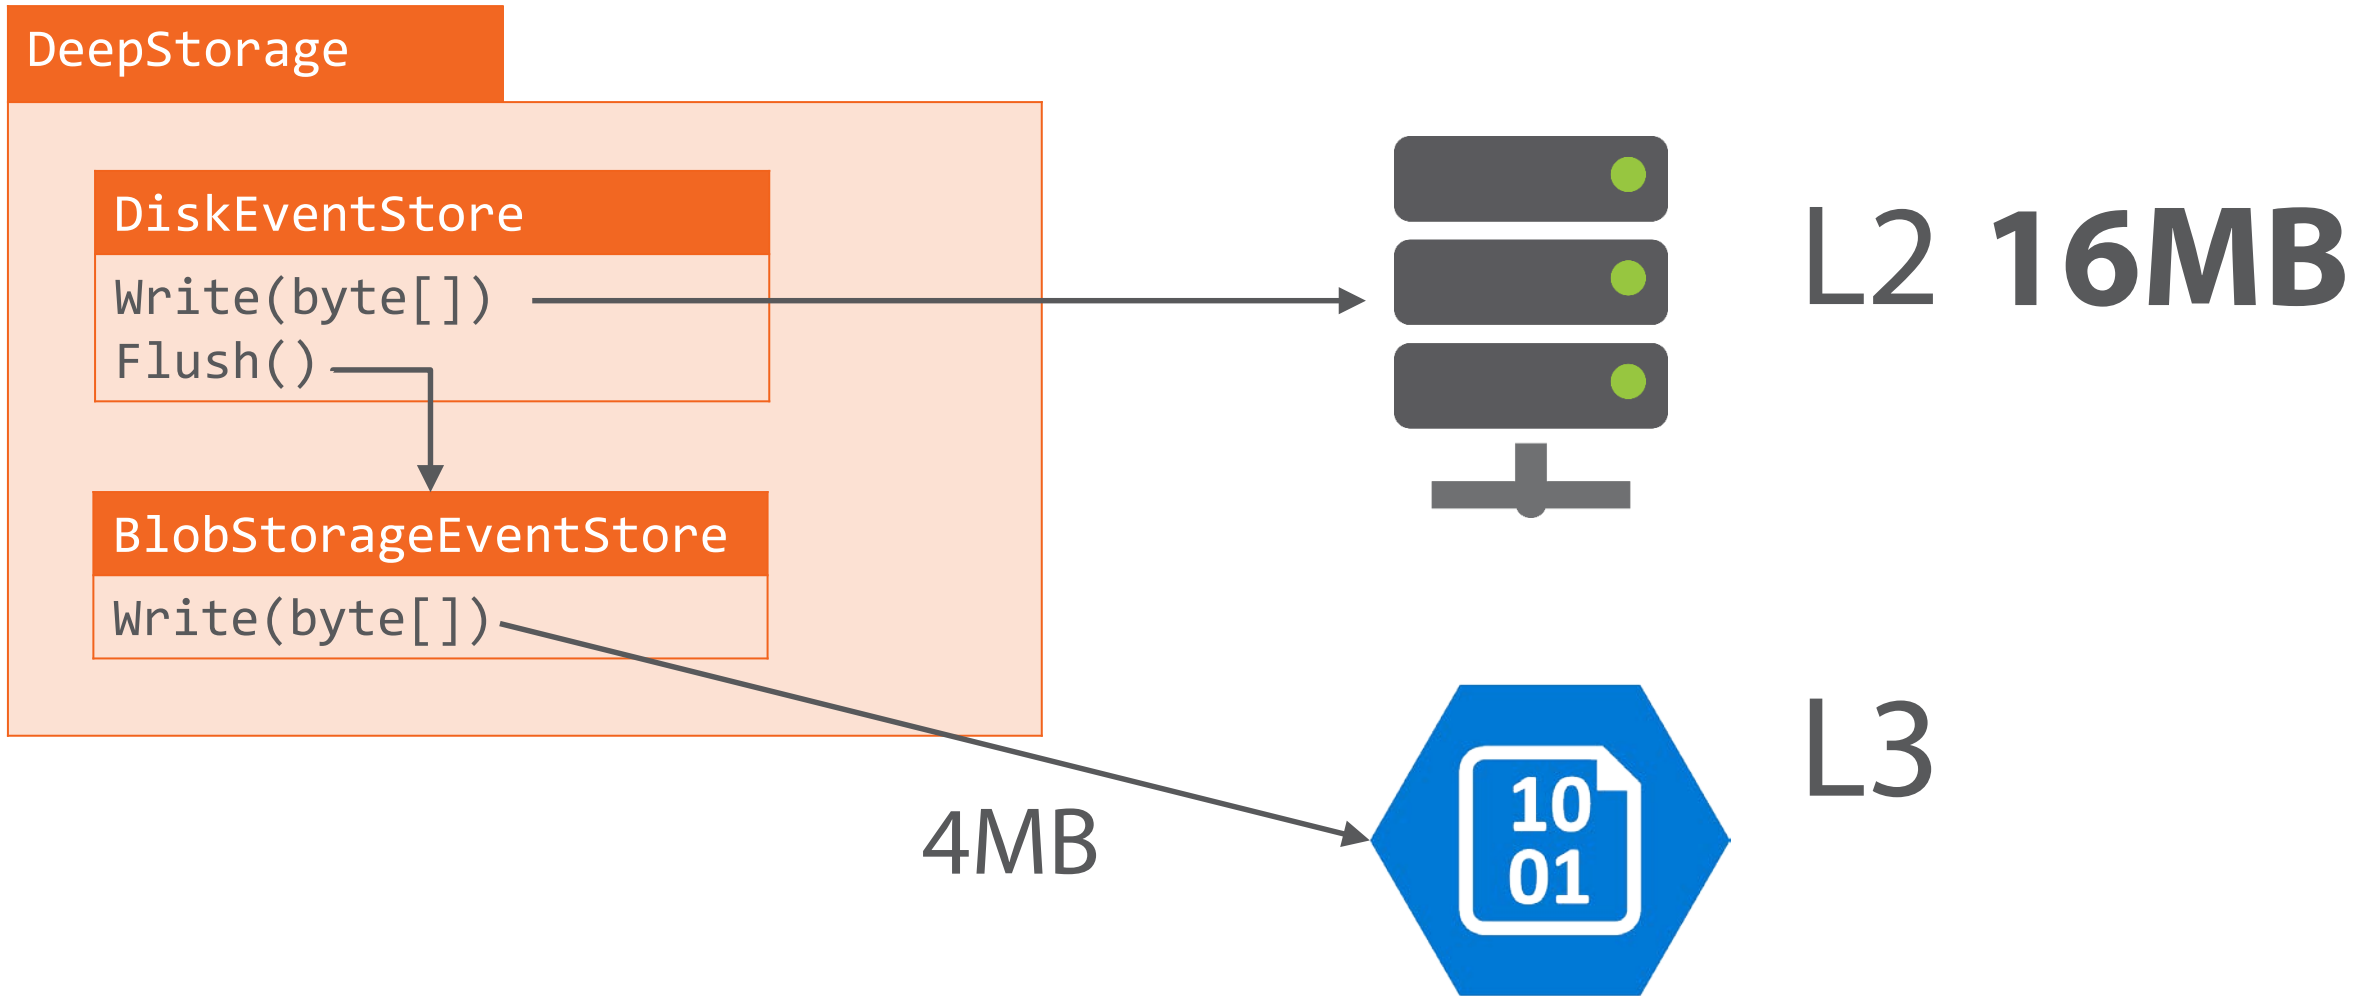
DeepStorage

DiskEventStore  
Write(byte[])



L2





## DeepStorage

### MemoryEventStore

Write(byte[])

Flush()


### DiskEventStore

Write(byte[])

Flush()

### BlobStorageEventStore

Write(byte[])



L1 8MB



L2 16MB



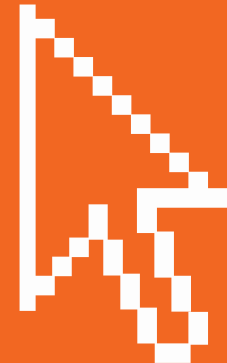
L3

# Demo: DiskEventStore

Compresses in memory

Appends GZip to disk

Inherits EventStoreBase



```
public override void Write(byte[] data)
{
    var compressedStream = new MemoryStream();
    using (var inputStream = new MemoryStream(data))
    {
```

## DiskEventStore

Write gets flushed events from MemoryEventStore



```
using (var compressionStream = new
GZipOutputStream(compressedStream))
{
    compressionStream.SetLevel(9);
    inputStream.CopyTo(compressionStream);
    compressionStream.Flush();
}
```

## DiskEventStore

Compress events in memory

```
_lock.Wait();  
using (var outputStream = File.OpenWrite(_filePath))  
{  
    outputStream.Position = outputStream.Length;  
    outputStream.Write(compressedData, 0, compressedData.Length);  
    outputStream.Flush();  
}
```

## DiskEventStore

Append compressed events to file

```
data = File.ReadAllBytes(_filePath);  
File.Delete(_filePath);  
using (File.Create(_filePath)) { }
```

## DiskEventStore

Flush reads & resets file

```
if (data.Length > 0)
{
    _startedTasks.Add(Task.Factory.StartNew(
        () => NextStore.Write(data)));
}
```

## DiskEventStore

Write to next store asynchronously

DeepStorage

BlobStorageEventStore

Write(byte[])



L3

DeepStorage

BlobStorageEventStore  
Write(byte[])



L3



DeepStorage

BlobStorageEventStore

Write(byte[])



L3



DeepStorage

BlobStorageEventStore

Write(byte[])



L3






DeepStorage

BlobStorageEventStore

```
Write(byte[])  
Flush() ←
```



L3



DeepStorage


BlobStorageEventStore

Write(byte[])

Flush()



L3



DeepStorage


BlobStorageEventStore

Write(byte[])

Flush()



L3



# Demo: BlobStorageEventStore

Inherits EventStoreBase

Appends  $\leq 4$ MB blocks

Commits with lease



```
_lock.Wait();  
using (var lease = new RenewingBlobLease(_blob))  
{  
    var offset = 0;  
    while (offset < data.Length)  
    {
```

## BlobStorageEventStore

Write loops data with renewing blob lease

```
var remaining = data.Length - offset;  
var length = remaining < MaxBlockSize ? remaining :  
                                                MaxBlockSize;  
using (var stream = new MemoryStream(data, offset, length))  
{
```

## BlobStorageEventStore

Extract blocks of 4MB or less

```
var blockId =  
    Convert.ToBase64String(Guid.NewGuid().ToArray());  
_blob.PutBlock(blockId, stream, null,  
    AccessCondition.GenerateLeaseCondition(lease.Id));  
offset += (int)stream.Length  
_blockIds.Add(blockId)
```

## BlobStorageEventStore


Append block with unique ID


```
var blockList =  
    _blob.DownloadBlockList(BlockListingFilter.Committed,  
        AccessCondition.GenerateLeaseCondition(lease.Id));  
var blockIds = blockList.Select(x => x.Name).ToList();  
blockIds.AddRange(_blockIds);  
_blob.PutBlockList(blockIds,  
    AccessCondition.GenerateLeaseCondition(lease.Id));
```


## BlobStorageEventStore

Flush commits new blocks









# Demo: WorkerRole

Start EventReceiver

Run WorkerRole

View blobs with CloudBerry Explorer



```
Container.Instance.RegisterType<IEventStore,  
                                MemoryEventStore>("1");  
Container.Instance.RegisterType<IEventStore,  
                                DiskEventStore>("2");  
Container.Instance.RegisterType<IEventStore,  
                                BlobStorageEventStore>("3");
```

## WorkerRole

Register event stores on Start

```
public override void Run()  
{  
    _receiver.RegisterProcessorAsync().Wait();  
    CompletedEvent.WaitOne();  
}
```


## WorkerRole

Register processor & start receiving events on Run

**16 partitions**

**2x small** instances

**16 per hour**



16 partitions



2x small instances



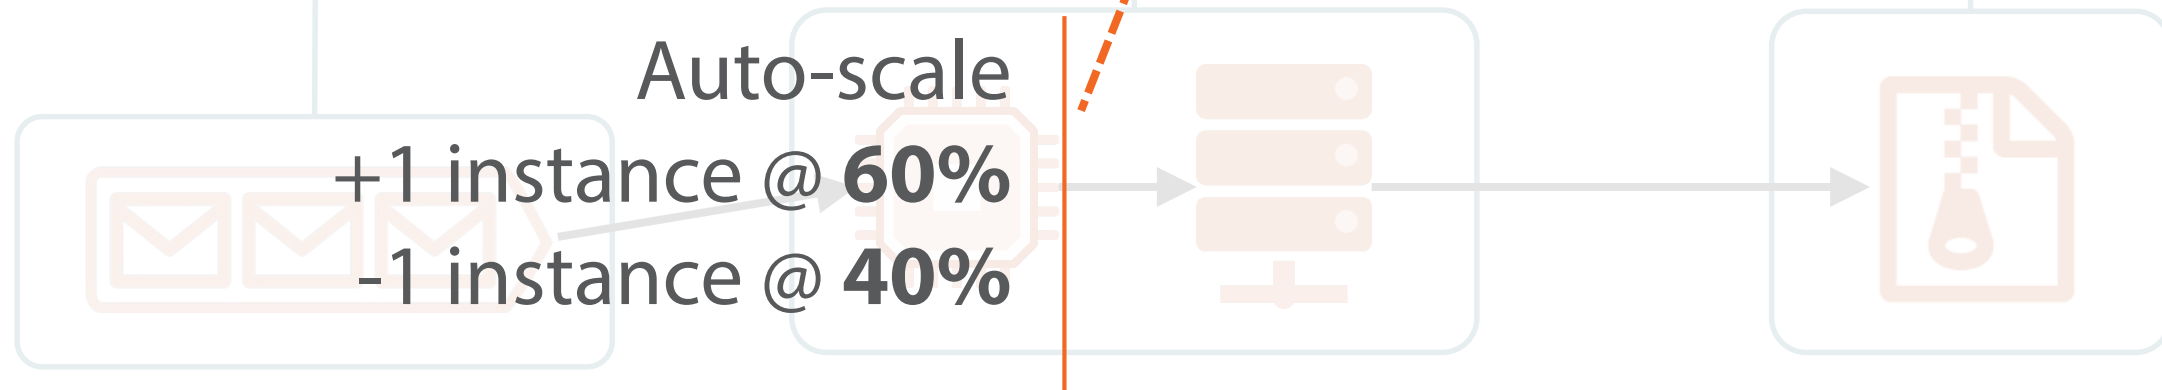
16 per hour



Auto-scale

+1 instance @ 60%

-1 instance @ 40%






**16 partitions**



**16x small instances**



**2.6K overnight**






1 throughput units  
**1MB/s** ingress



**1** throughput units  
**1MB/s** ingress  
**2MB/s** egress




**5** throughput units  
**5MB/s** ingress  
**10MB/s** egress



- 500TB storage
- 20K requests /sec
- 20Gb/sec upload
- 30Gb/sec download




Per Storage Account



- 500TB storage
- 20K requests /sec
- 20Gb/sec upload
- 30Gb/sec download




**<\$25** per TB per month



## EventProcessorHost


- Receive events
- Lock partitions
- Checkpoint progress

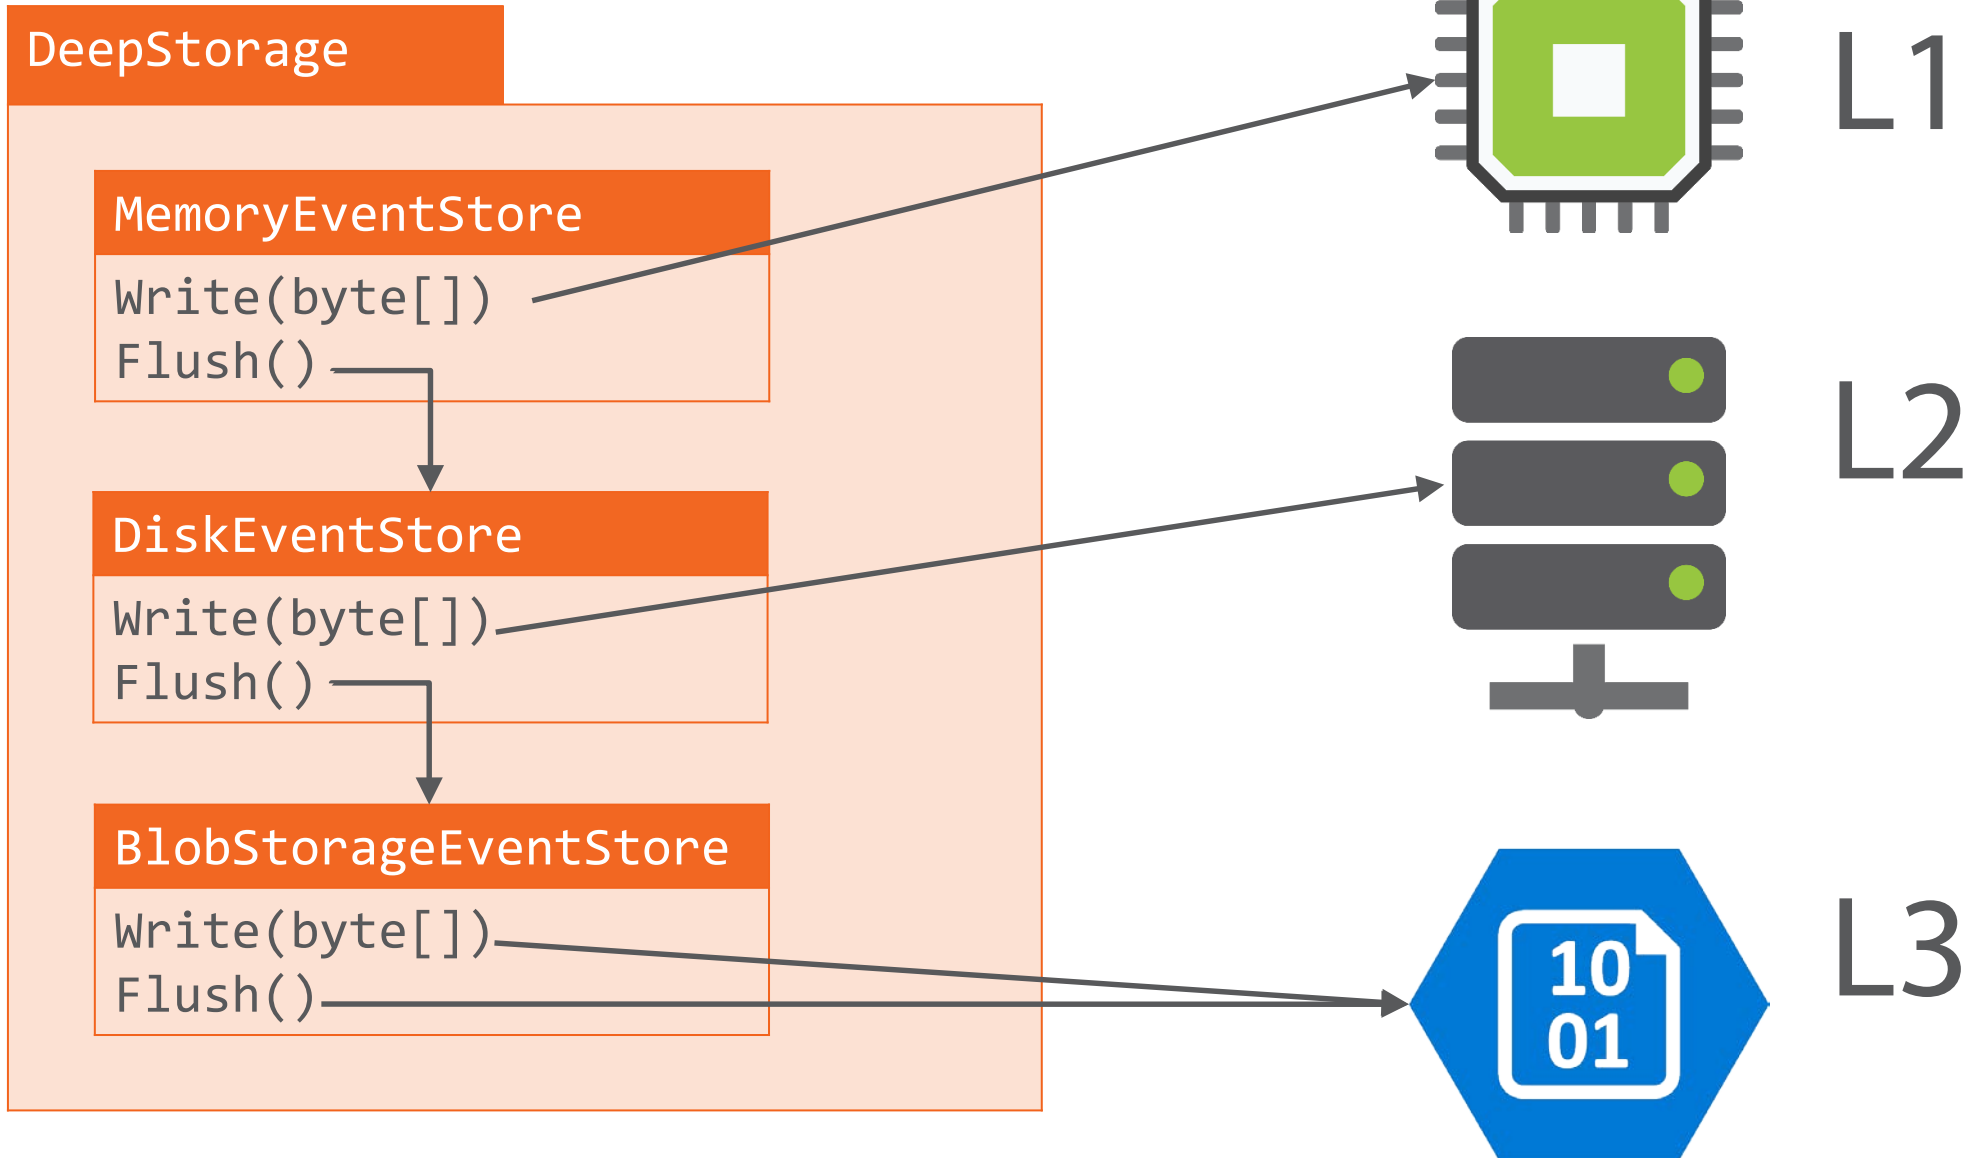


**`{partition}/{hour}.json.gz`**

- GZipped JSON
- One partition, one hour
- 384 files per day









p1/2015033101.json.gz

p1/2015033102.json.gz

...

p15/2015040122.json.gz

p15/2015040123.json.gz

