


# Normalizing Event Data for Real-Time Queries





Elton Stoneman

@EltonStoneman | [www.geekswithblogs.net/eltonstoneman](http://www.geekswithblogs.net/eltonstoneman)

---









p1/2015033101.json.gz

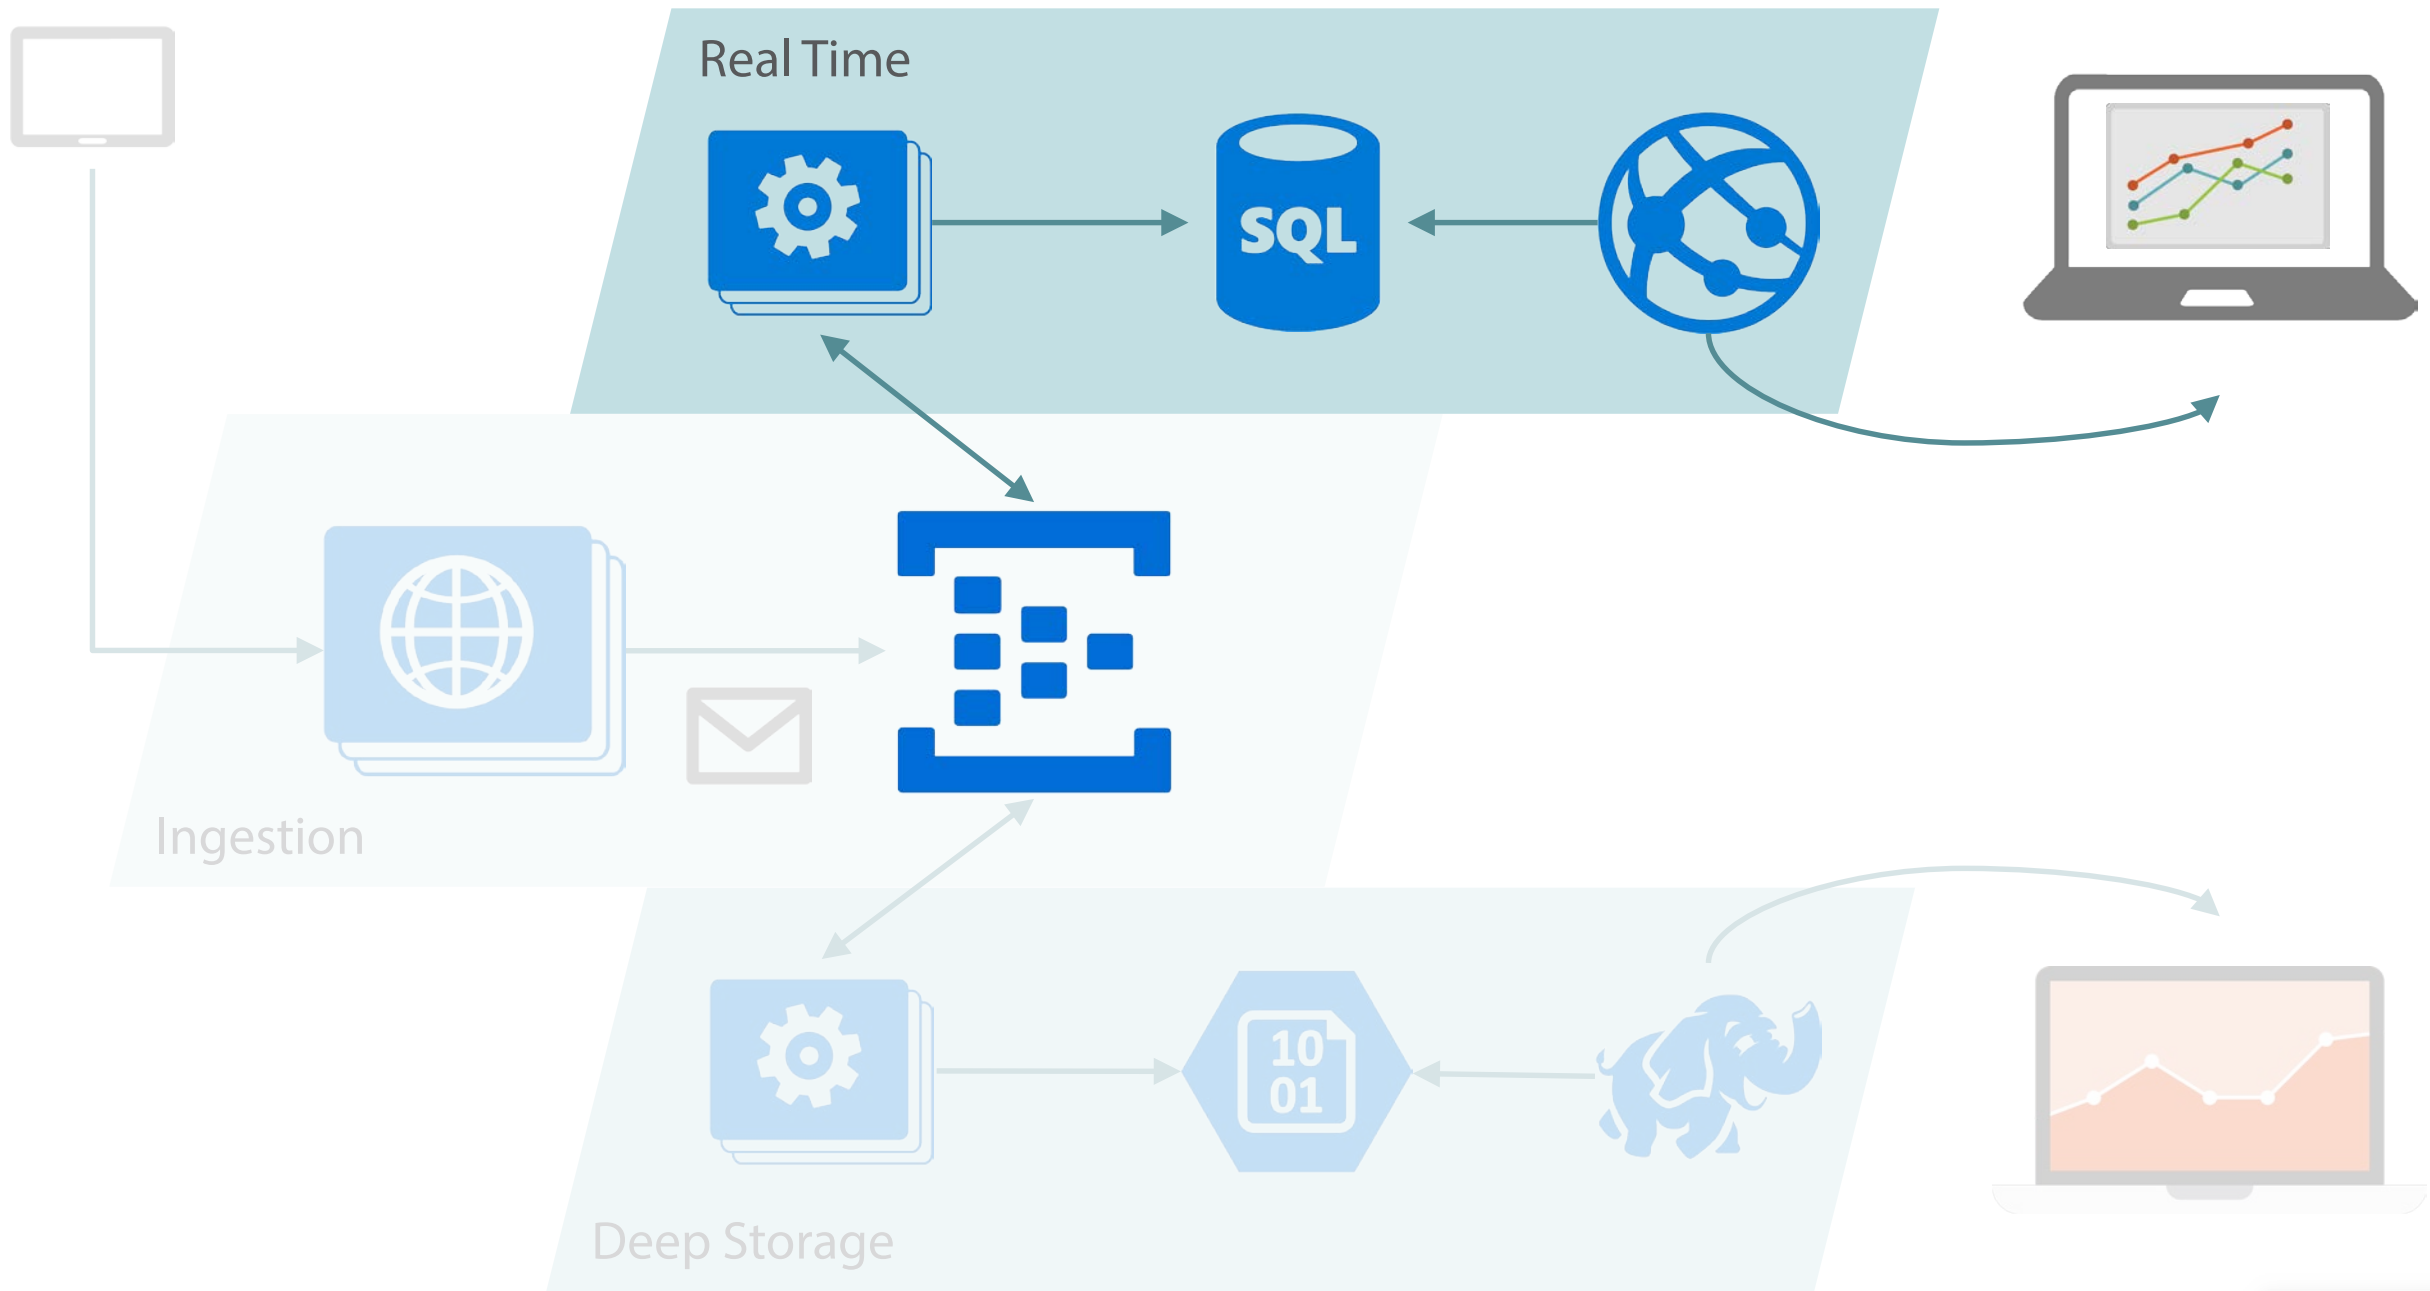
p1/2015033102.json.gz

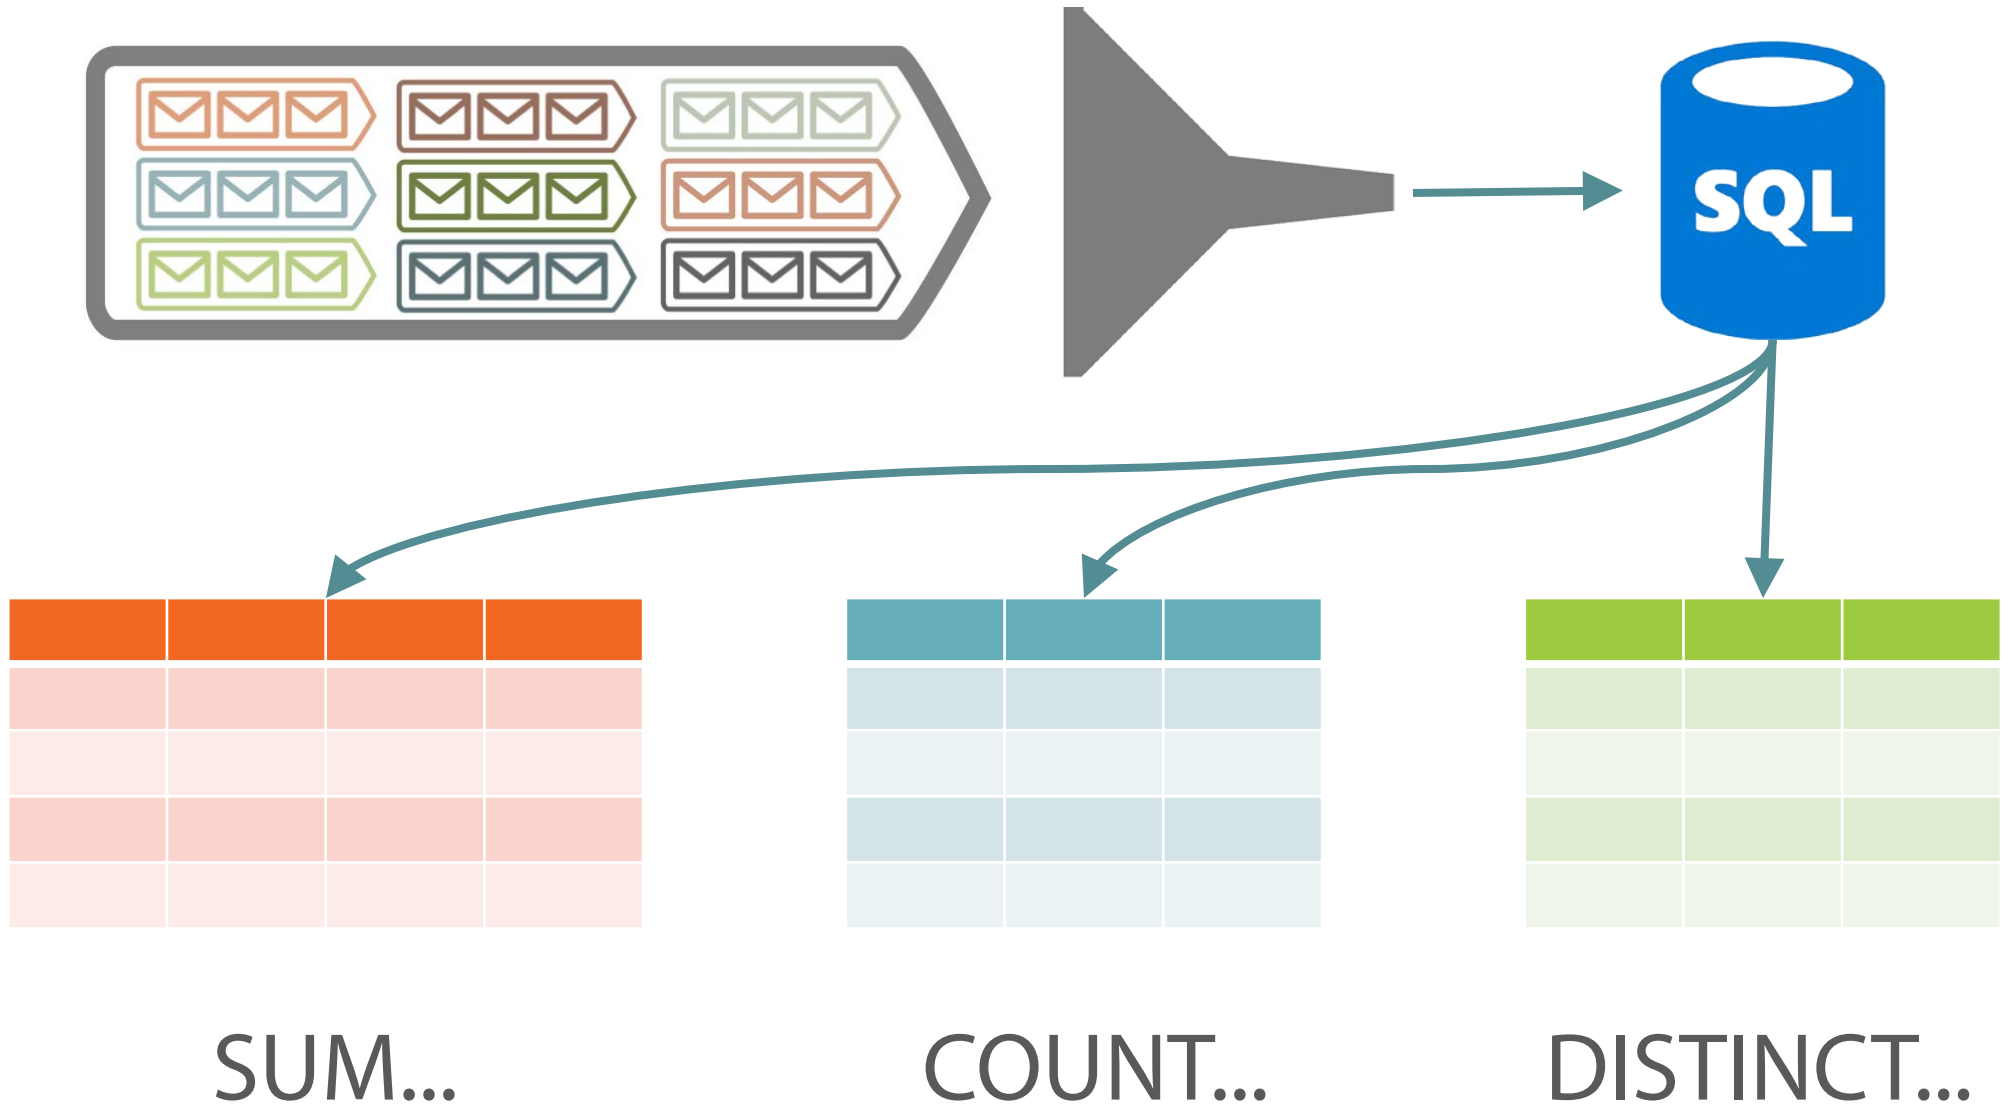
...

p15/2015040122.json.gz

p15/2015040123.json.gz













SUM...

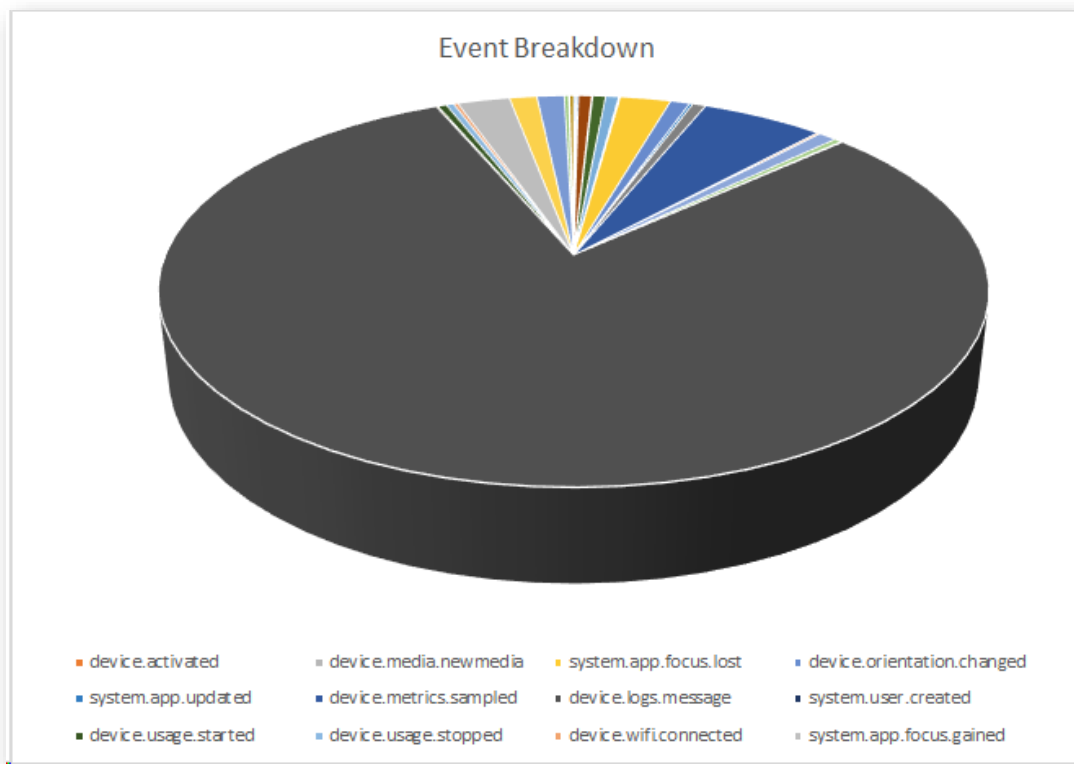

Filter


Aggregate


COUNT...


DISTINCT...








Real Time


Today's Events  
**506.8M**

Last updated at 22:09.27



gained  
lost  
updated  
installed  
removed

Last updated at 22:09.28



switched  
updated  
created  
deleted

Last updated at 22:09.28




Ingest

### Top Events


device.logs.message	409.8M
device.metrics.sampled	28.0M
system.app.focus.gained	12.6M
system.app.focus.lost	12.2M
device.gps.usage.stopped	5.9M
(other)	38.4M

Last updated at 22:09.28



logs.message  
metrics.sampled  
app.focus.gained  
app.focus.lost  
gps.usage.stopped  
(other)

Last updated at 22:09.28



connected  
disconnected  
enabled  
disabled

Last updated at 22:09.28

Deep Storage

Real Time





## EventMetrics

Period (NVARCHAR

EventName (NVARCHAR)

Count (BIGINT)

Period	EventName	Count
2015050617	device.logs.message	3,336,902
2015050617	device.metrics.sent	124,9426
2015050618	device.logs.message	4,206,7139
2015050618	device.metrics.sent	213, 807



### EventMetrics

Period (NVARCHAR

EventName (NVARCHAR)

Count (BIGINT)

@40 event types  
**960 rows /day**

Period	EventName	Count
2015050617	device.logs.message	3,336,902
2015050617	device.metrics.sent	124,9426
2015050618	device.logs.message	4,206,7139
2015050618	device.metrics.sent	213, 807





## EventMetrics


Period (NVARCHAR


EventName (NVARCHAR)

Count (BIGINT)

Period	EventName	Count
2015050617	device.logs.message	3,336,902
2015050617	device.metrics.sent	124,926
2015050618	device.logs.message	4,206,739
2015050618	device.metrics.sent	213,807









## EventMetrics

Period (NVARCHAR)

**PartitionId (NVARCHAR)**

EventName (NVARCHAR)

Count (BIGINT)

Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	336,902
2015050617	5	device.logs.message	324,426
2015050617	9	device.logs.message	306,139
2015050617	11	device.logs.message	313,807



## EventMetrics

Period (NVARCHAR

**PartitionId (NVARCHAR)**

EventName (NVARCHAR)

Count (BIGINT)

@40 event types

@16 partitions

**15,360 rows  
/day**

Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	336,902
2015050617	5	device.logs.message	324,426
2015050617	9	device.logs.message	306,139
2015050617	11	device.logs.message	313,807



## EventMetrics

Period (NVARCHAR

PartitionId (NVARCHAR)

EventName (NVARCHAR)

Count (BIGINT)

IX\_EventMetrics




Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	336,902
2015050617	5	device.logs.message	324,426
2015050617	9	device.logs.message	306,139
2015050617	11	device.logs.message	313,807

# Demo: Events DB

SQL Database Project

Sample data

Entity Framework model

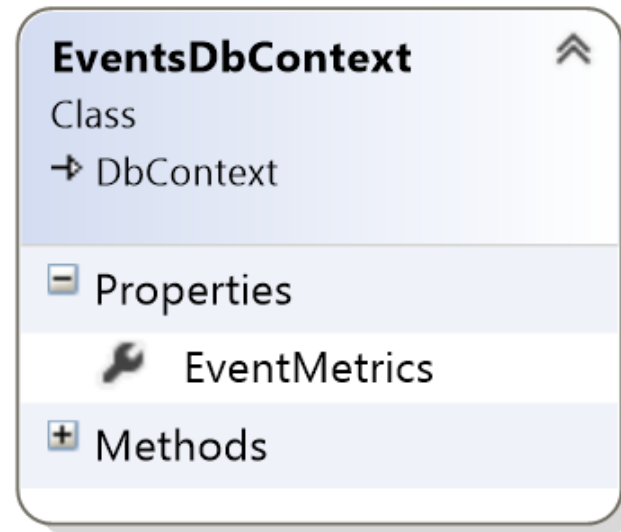
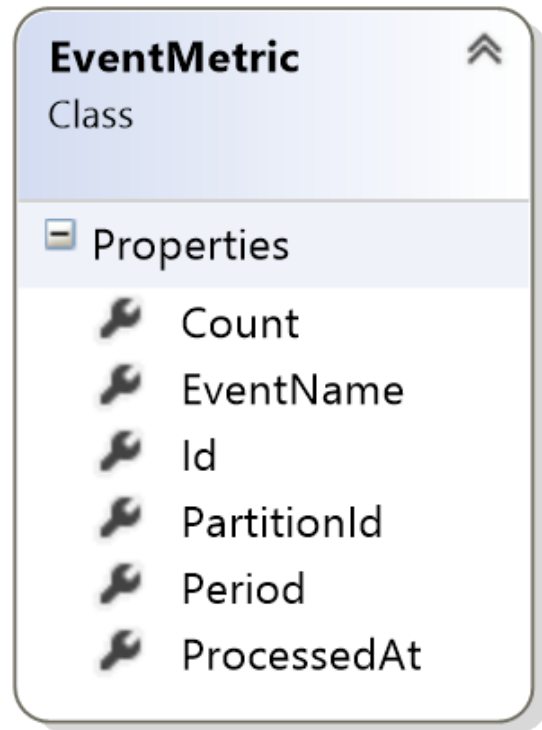


```
CREATE TABLE [dbo].[EventMetrics] (  
  [Id] INT PRIMARY KEY IDENTITY,  
  [EventName] NVARCHAR(40),  
  [Period] NVARCHAR(10), [PartitionId] NVARCHAR(5),  
  [Count] BIGINT, [ProcessedAt] DATETIME )
```

## EventMetrics


Aggregated event data and processing timestamp






# Entity Framework

Model & database context





Real Time





Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426
2015050617	9	device.logs.message	36,139
2015050617	11	device.logs.message	33,807



SELECT \* FROM EventMetrics



Period	PartitionId	EventName	Count
2015050617	13	<u>device.logs.message</u>	36,902
2015050617	5	<u>device.logs.message</u>	34,426
2015050617	9	<u>device.logs.message</u>	36,139
2015050617	11	<u>device.logs.message</u>	33,807




INSERT INTO EventMetrics




Period	PartitionId	EventName	Count
2015050617	13	<u>device.logs.message</u>	36,902
2015050617	5	<u>device.logs.message</u>	34,426
2015050617	9	<u>device.logs.message</u>	36,139
2015050617	11	<u>device.logs.message</u>	33,807



UPDATE EventMetrics




Period	PartitionId	EventName	Count
2015050617	13	<u>device.logs.message</u>	36,902
2015050617	5	<u>device.logs.message</u>	34,426
2015050617	9	<u>device.logs.message</u>	36,139
2015050617	11	<u>device.logs.message</u>	33,807



Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426
2015050617	9	device.logs.message	36,139
2015050617	11	device.logs.message	33,807






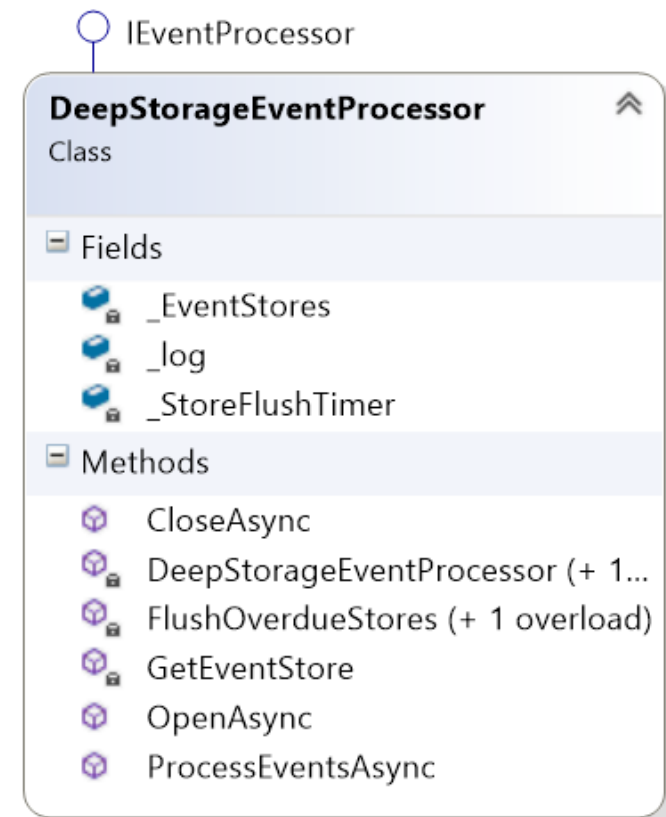
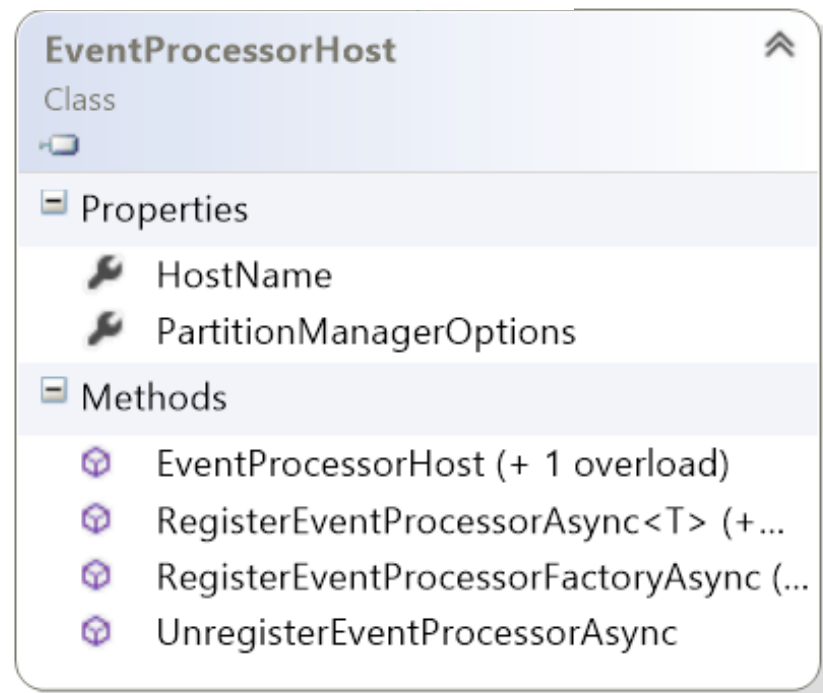
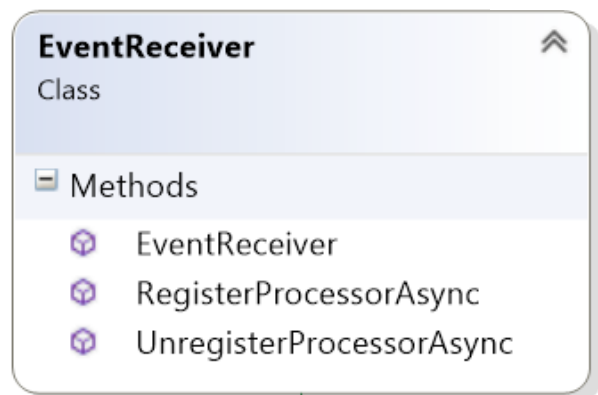


Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426
2015050617	9	device.logs.message	36,139
2015050617	11	device.logs.message	33,807



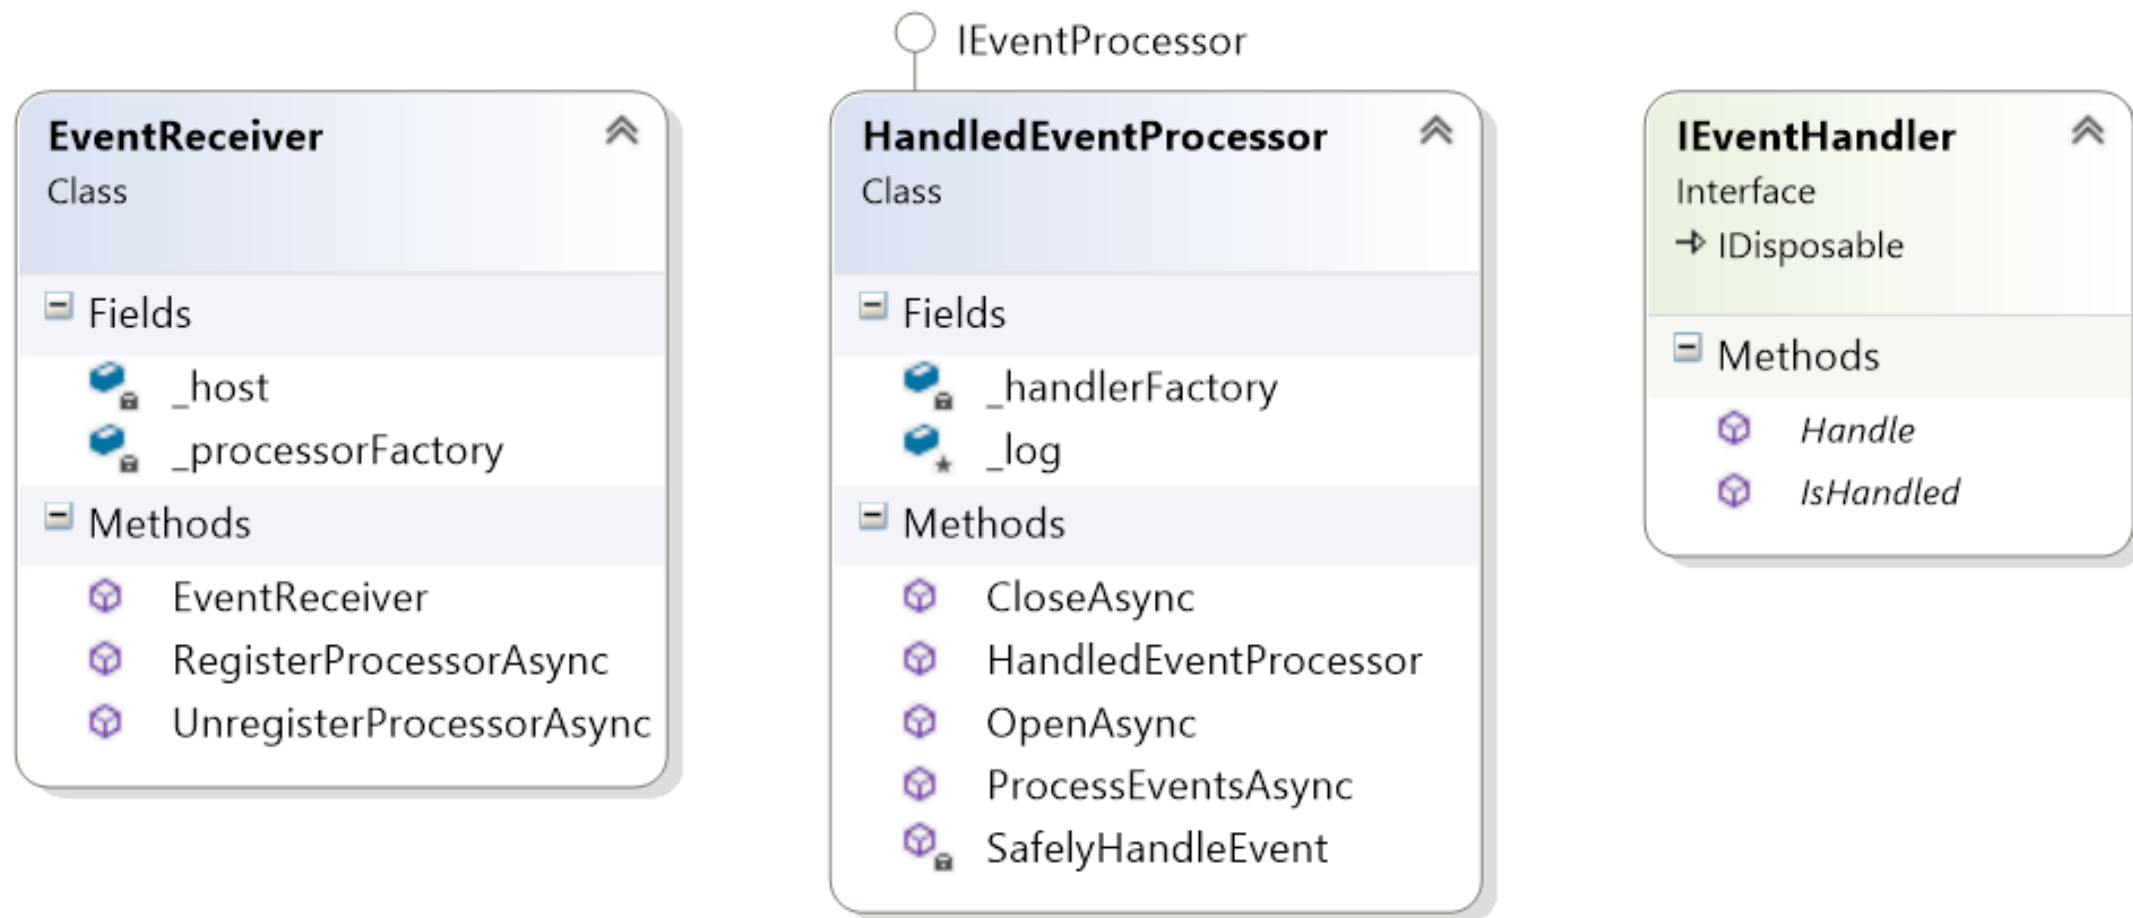


		Period	PartitionId	EventName	Count
		2015050617	13	device.logs.message	36,902
		2015050617	5	device.logs.message	34,426
		2015050617	9	device.logs.message	36,139
		2015050617	11	device.logs.message	33,807






# Deep Storage Worker

EventReceiver, EventProcessorHost and IEventProcessor



## Refactored Deep Storage Worker

EventReceiver, EventProcessorHost and IEventProcessor













RealTime.Worker

EventMetricsHandler

FirmwareVersionHandler




# Demo: Core.EventProcessor

Refactored Deep Storage Worker

Common EventReceiver

Custom IEventHandler



```
public interface IEventHandler : IDisposable
{
    bool IsHandled(string eventName);
    void Handle(EventData eventData, string partitionId);
}
```

## IEventHandler

Operates on single EventData

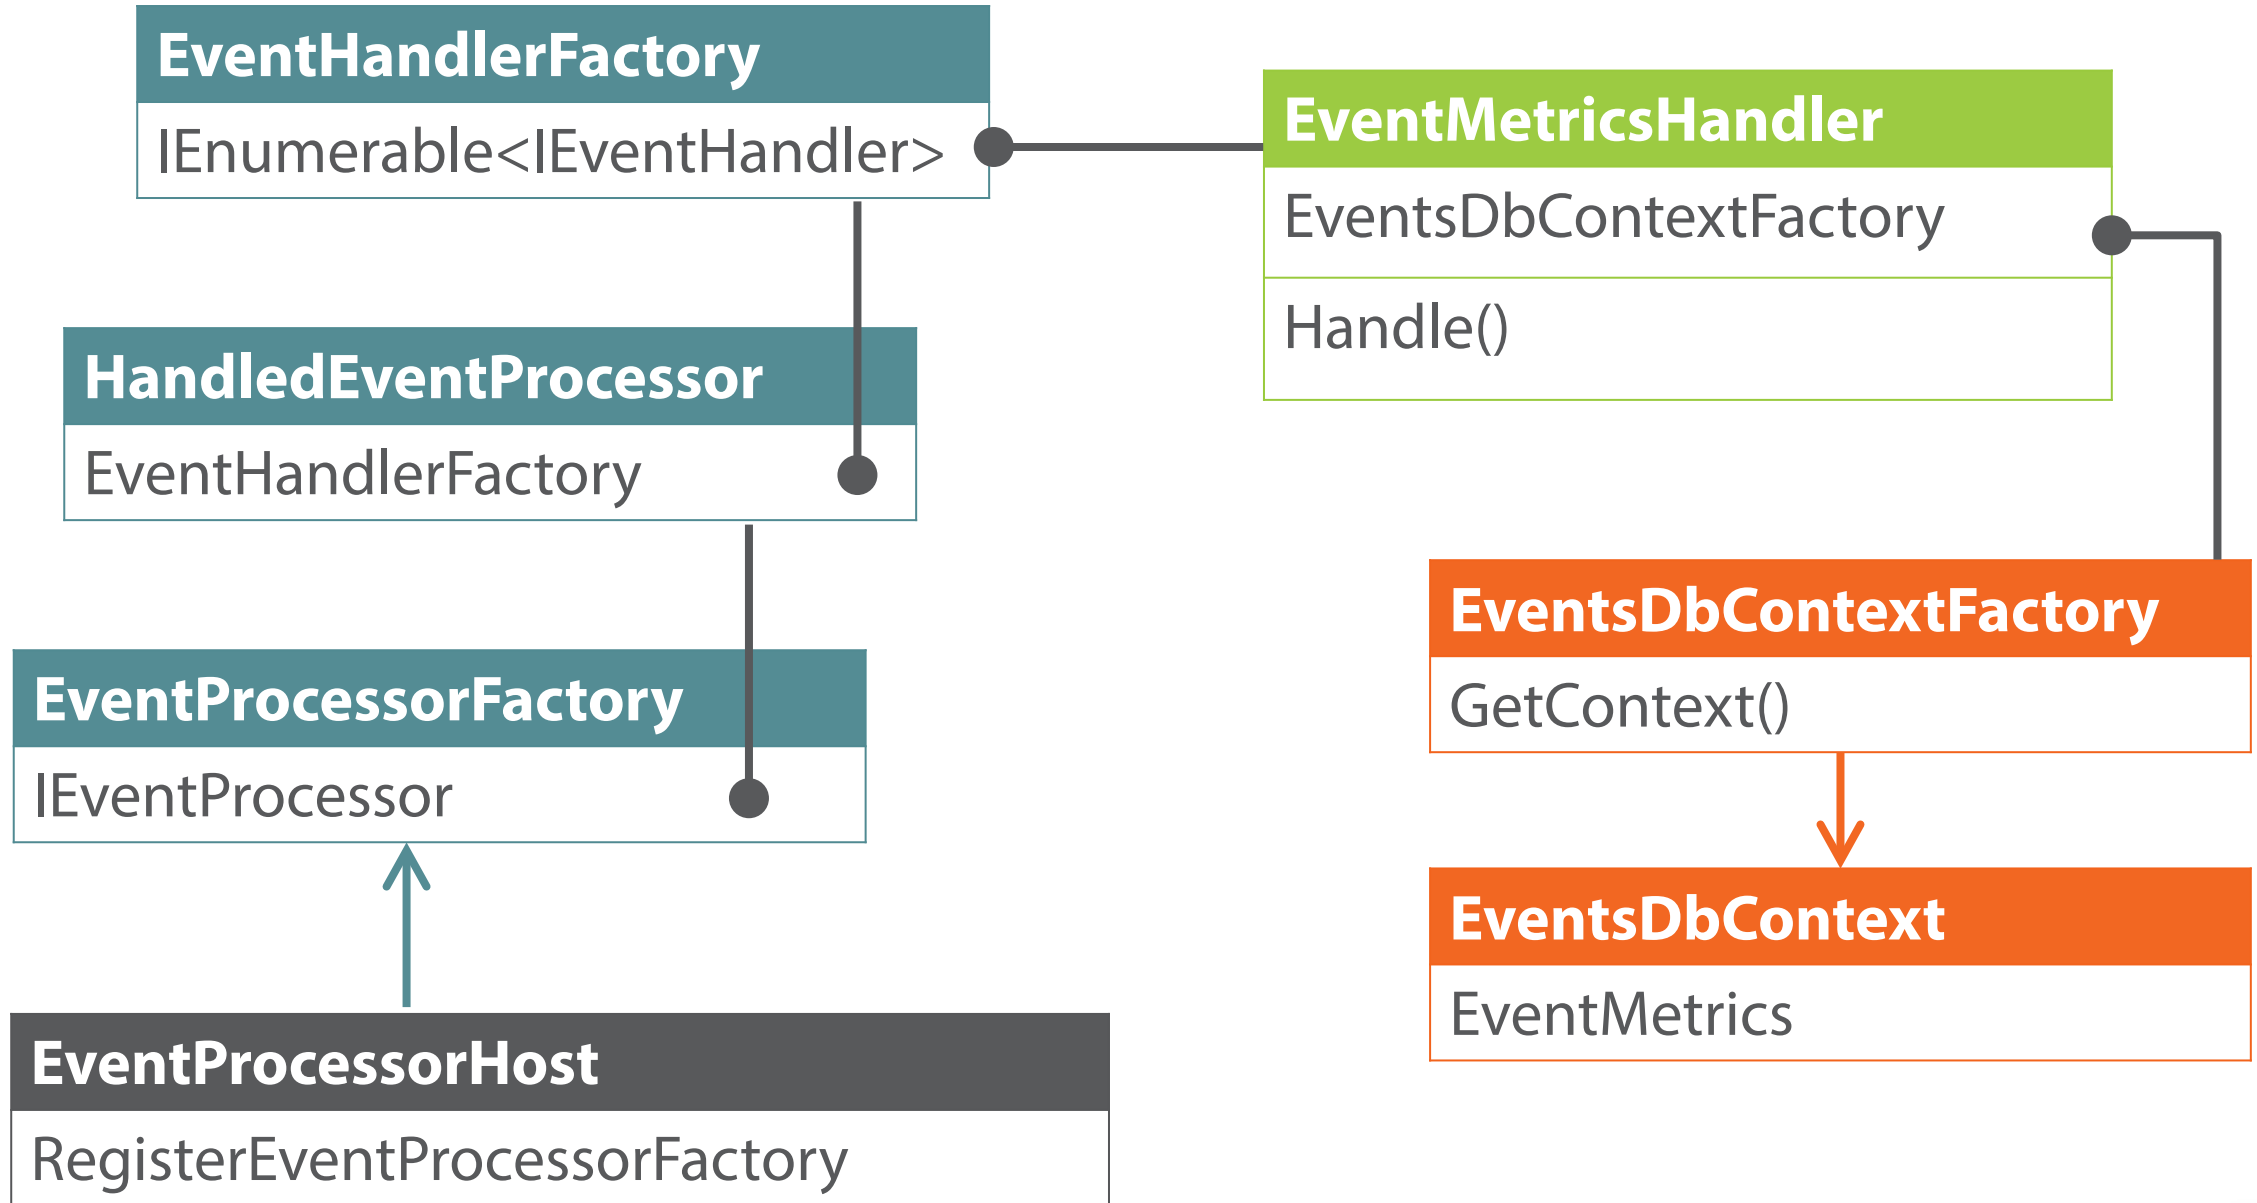




```
{  
  "timestamp": 1415878188000,  
  "message": "Unhandled exception",  
  "eventName": "device.logs.message",  
  "deviceId": "1e565484571752"  
}
```

eventName	device.logs.message
receivedAt	14158792533000

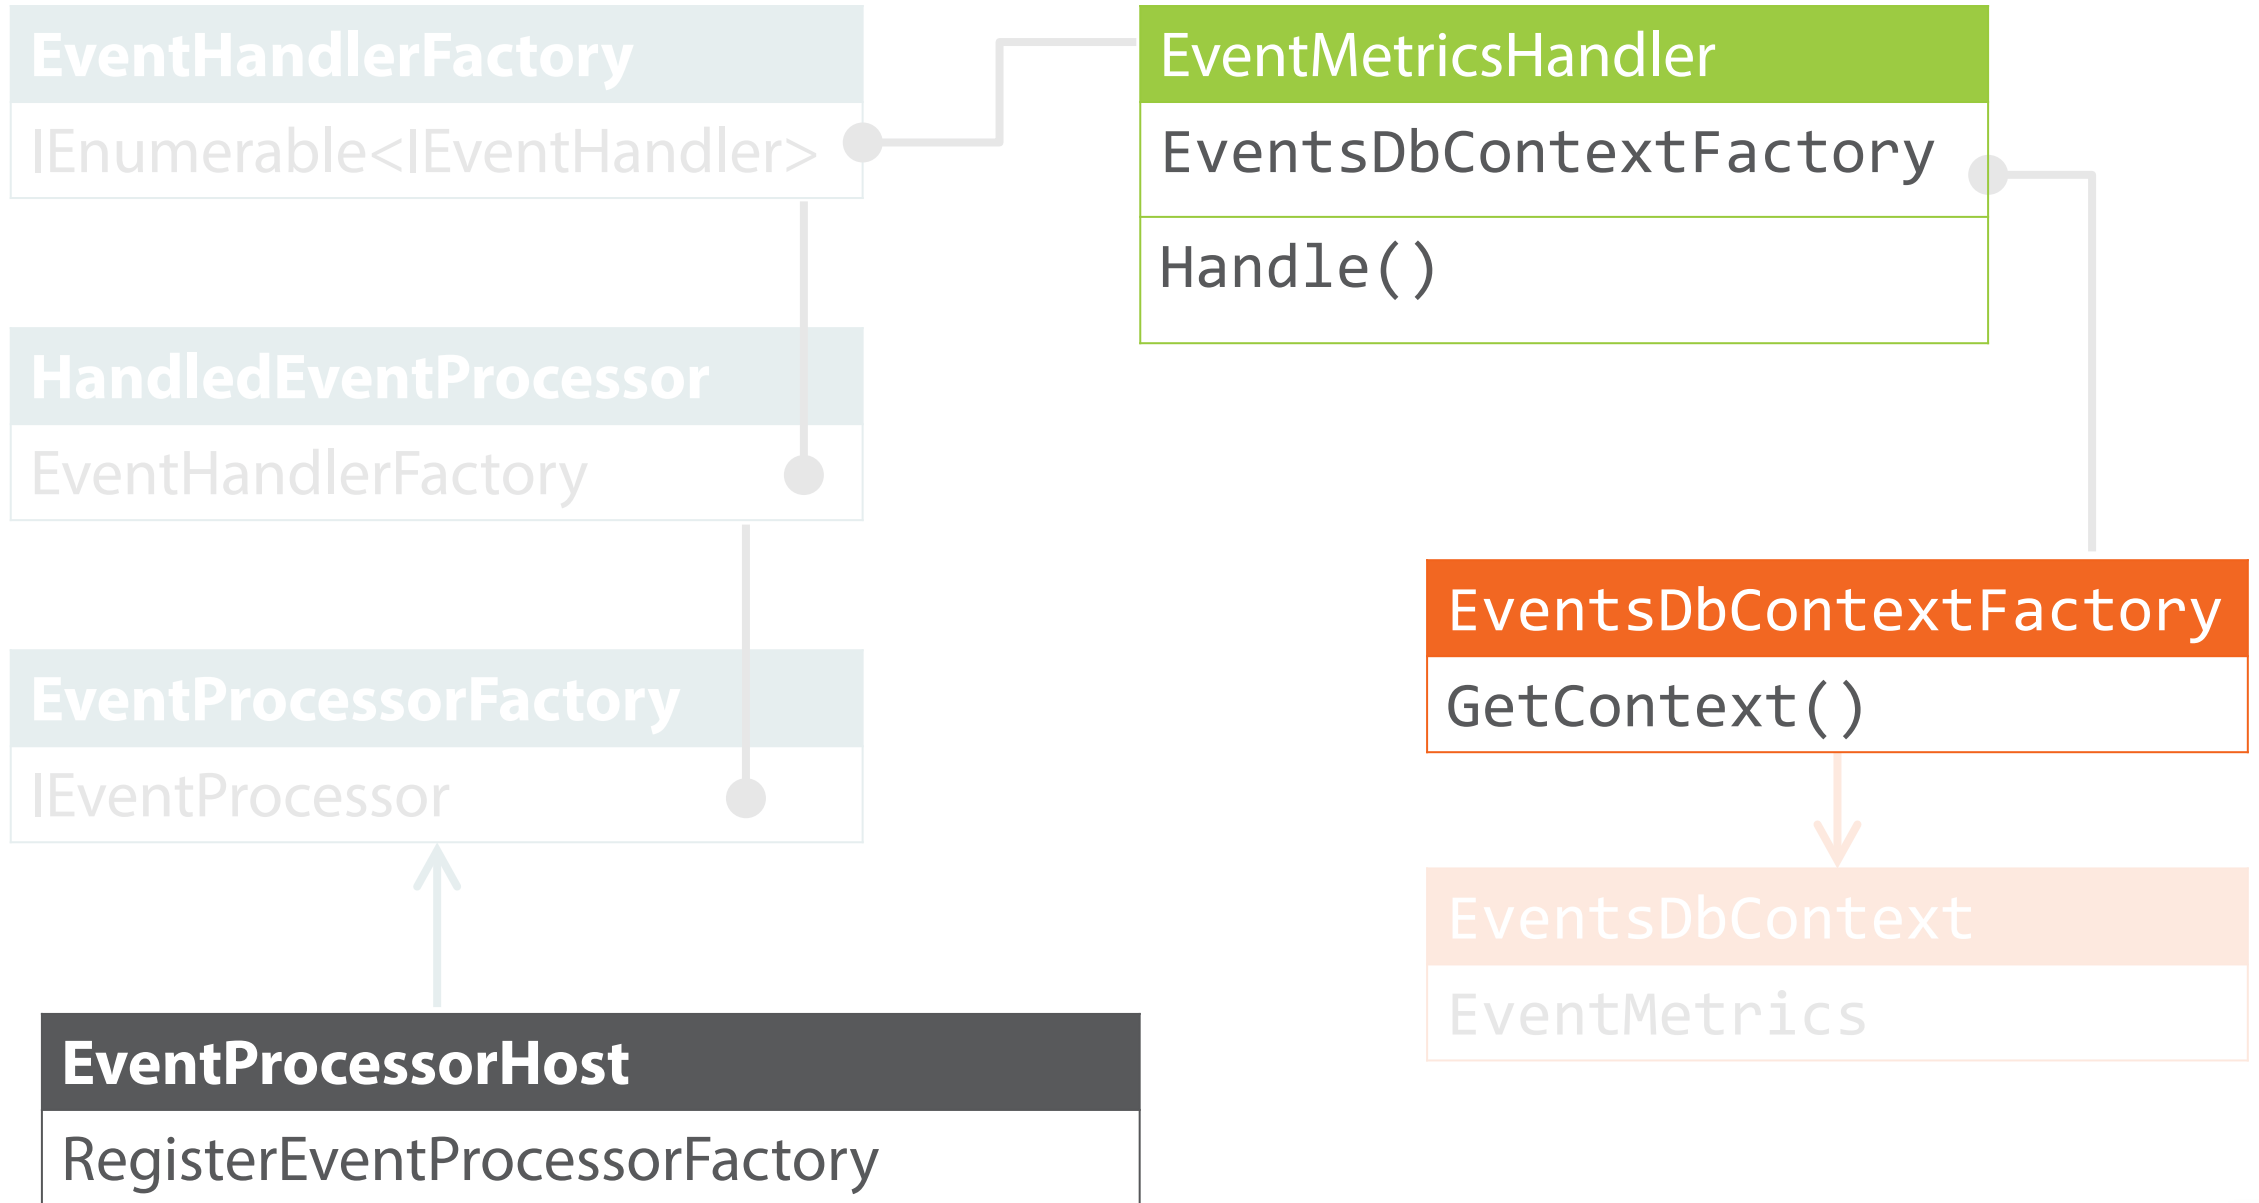
partitionId	11
-------------	----



```
container.RegisterType<IEventHandler,  
    DeepStorageEventHandler>("DeepStorageEventHandler");  
container.RegisterType<EventHandlerFactory>();  
  
container.RegisterType<IEventProcessor, HandledEventProcessor>();  
container.RegisterType<IEventProcessorFactory, EventProcessorFactory>();
```

## Worker Role

Register dependency chain



**RealTime**

**EventMetricsHandler**

Handle(EventData)



eventName

device.logs.message

receivedAt

14158792533000

Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426


**RealTime**

**EventMetricsHandler**

Handle(EventData)



Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426




Period	PartitionId	EventName	Count
2015050617	13	device.logs.message	36,902
2015050617	5	device.logs.message	34,426



**RealTime**

**EventMetricsHandler**

Handle(EventData)



```
{  
  "period":2015050717,  
  "partitionId" : "11",  
  "eventName": "event.a",  
  "count": 32016  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "6",  
  "eventName": "event.a",  
  "count": 31943  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "1",  
  "eventName": "event.b",  
  "count": 184  
}
```



**RealTime**

**EventMetricsHandler**

Handle(EventData)



eventName

event.a

receivedAt

14158792533000

partitionId

6

```
{
  "period":2015050717,
  "partitionId" : "11",
  "eventName": "event.a",
  "count": 32016
}
```

```
{
  "period":2015050717,
  "partitionId" : "6",
  "eventName": "event.a",
  "count": 31943
}
```

RealTime

EventMetricsHandler

Handle(EventData)

++

```
{  
  "period":2015050717,  
  "partitionId" : "11",  
  "eventName": "event.a",  
  "count": 32016  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "6",  
  "eventName": "event.a",  
  "count": 31944  
}
```



eventName	event.a
receivedAt	14158792533000
partitionId	6

**RealTime**

**EventMetricsHandler**

Handle(EventData)



eventName	event.b
receivedAt	14158792514000
partitionId	1

```
{
  "period":2015050717,
  "partitionId" : "11",
  "eventName": "event.a",
  "count": 32016
}
```

```
{
  "period":2015050717,
  "partitionId" : "6",
  "eventName": "event.a",
  "count": 31944
}
```

**RealTime**

**EventMetricsHandler**

Handle(EventData)




eventName	event.b
receivedAt	14158792514000
partitionId	1

```
{  
  "period":2015050717,  
  "partitionId" : "11",  
  "eventName": "event.a",  
  "count": 32016  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "6",  
  "eventName": "event.a",  
  "count": 31943  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "1",  
  "eventName": "event.b",  
  "count": 1  
}
```



eventName	event.b
receivedAt	14158792514000
partitionId	1

`ConcurrentDictionary<string, EventMetric>`

```
{  
  "period":2015050717,  
  "partitionId" : "11",  
  "eventName": "event.a",  
  "count": 32016  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "6",  
  "eventName": "event.a",  
  "count": 31943  
}
```

```
{  
  "period":2015050717,  
  "partitionId" : "1",  
  "eventName": "event.b",  
  "count": 184  
}
```


## RealTime

### ModelBuffer<T>

ConcurrentDictionary<string, T>

Add(T)

Flush(IEnumerable<T>)




# Demo: ModelBuffer<T>

Buffer in-memory

Flush at elapsed time

Used by Event Handler



```
public class ModelBuffer<TModel>
{
    private Action<IEnumerable<TModel>> _doSave;
    private Func<TModel, string> _getKey;
    private ConcurrentDictionary<string, TModel> _models;
```

ModelBuffer<TModel>

TModel is a code-first EF type



```
_models.TryAdd(key, model)  
_models.TryGetValue(key, out model)  
_models.Select(x => x.Value)
```

## ConcurrentDictionary<TKey, TValue>

Use threadsafe methods with minimal locking

```
var buffer = new ModelBuffer<StubModel>(
    x => SaveModels(x),
    y => string.Format("{0}p{1}i{2}",
                        y.Period, y.PartitionId, y.Index),
    TimeSpan.FromMilliseconds(200));
```

## ModelBuffer<T> usage

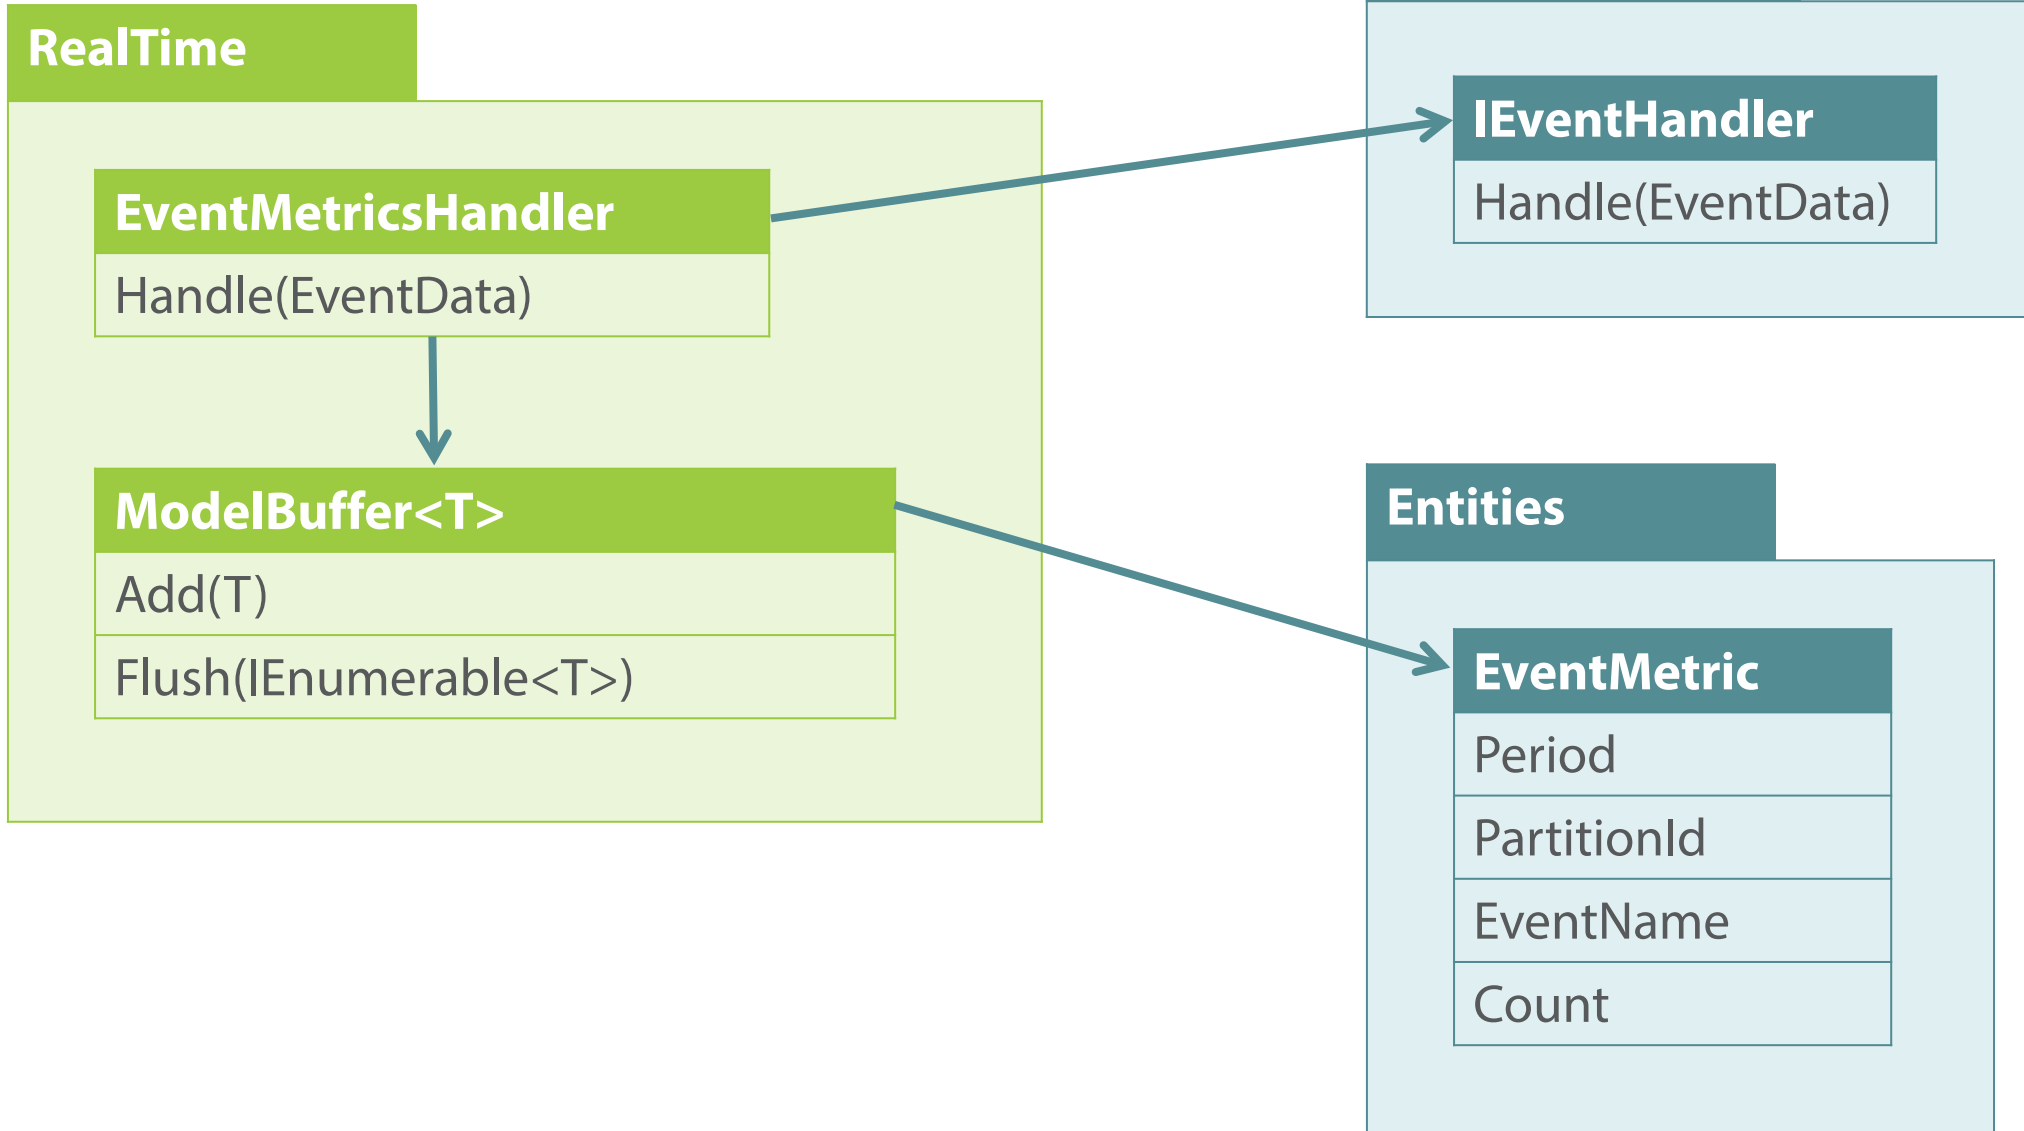
Flush action, key builder func and flush period


```
<add key="ModelBuffers.EventMetric.BufferFlushTime"  
      value="00:00:00.05" />
```

```
<add key="ModelBuffers.BufferFlushTime"  
      value="00:00:05" />
```

## App Settings

Custom & default flush periods






```
{  
  "period":2015050720,  
  "eventName": "event.a",  
  "count": 32016  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.b",  
  "count": 12961  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.c",  
  "count": 184  
}
```



```
{  
  "period":2015050720,  
  "eventName": "event.a",  
  "count": 51974  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.b",  
  "count": 18621  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.c",  
  "count": 348  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.a",  
  "count": 1574  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.b",  
  "count": 357  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.c",  
  "count": 2  
}
```



```
{  
  "period":2015050720,  
  "eventName": "event.a",  
  "count": 51974  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.b",  
  "count": 18621  
}
```

```
{  
  "period":2015050720,  
  "eventName": "event.c",  
  "count": 348  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.a",  
  "count": 1574  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.b",  
  "count": 357  
}
```

```
{  
  "period":2015050721,  
  "eventName": "event.c",  
  "count": 2  
}
```



Period	PartitionId	EventName	Count
2015050617	1	device.logs.message	336,902
2015050617	2	device.logs.message	324,426
...			
2015050617	14	device.logs.message	306,139
2015050617	15	device.logs.message	313,807




# Demo: EventMetricsHandler

Implements IEventHandler

DI for EventsDbContext

ModelBuffer<EventMetric>



```
public class EventMetricsHandler : IEventHandler
{
    public bool IsHandled(string eventName)
    {
        return true;
    }
}
```

## EventMetricsHandler

Handles all event types

```
var eventName = eventData.GetEventName();  
var period = eventData.GetReceivedAtHour();  
var metric = GetMetric(eventName, period, partitionId);  
metric.Count += 1;  
metric.ProcessedAt = DateTime.UtcNow;
```


## Handle(EventData)

Finds EventMetric from buffer and updates

```
var existingMetric = db.EventMetrics.FirstOrDefault(  
    x => x.PartitionId == metric.PartitionId &&  
        x.Period == metric.Period &&  
        x.EventName == metric.EventName);
```

## SaveMetrics

Checks for row in database and updates/inserts








# Demo: Real Time Worker Role

Configure container

Start EventReceiver

Update SQL Azure





```
container.RegisterType<IEventHandler,  
    EventMetricsHandler>("EventMetricsHandler");  
container.RegisterType<EventHandlerFactory>();  
container.RegisterType<IEventProcessor, HandledEventProcessor>();  
container.RegisterType<IEventProcessorFactory, EventProcessorFactory>();
```



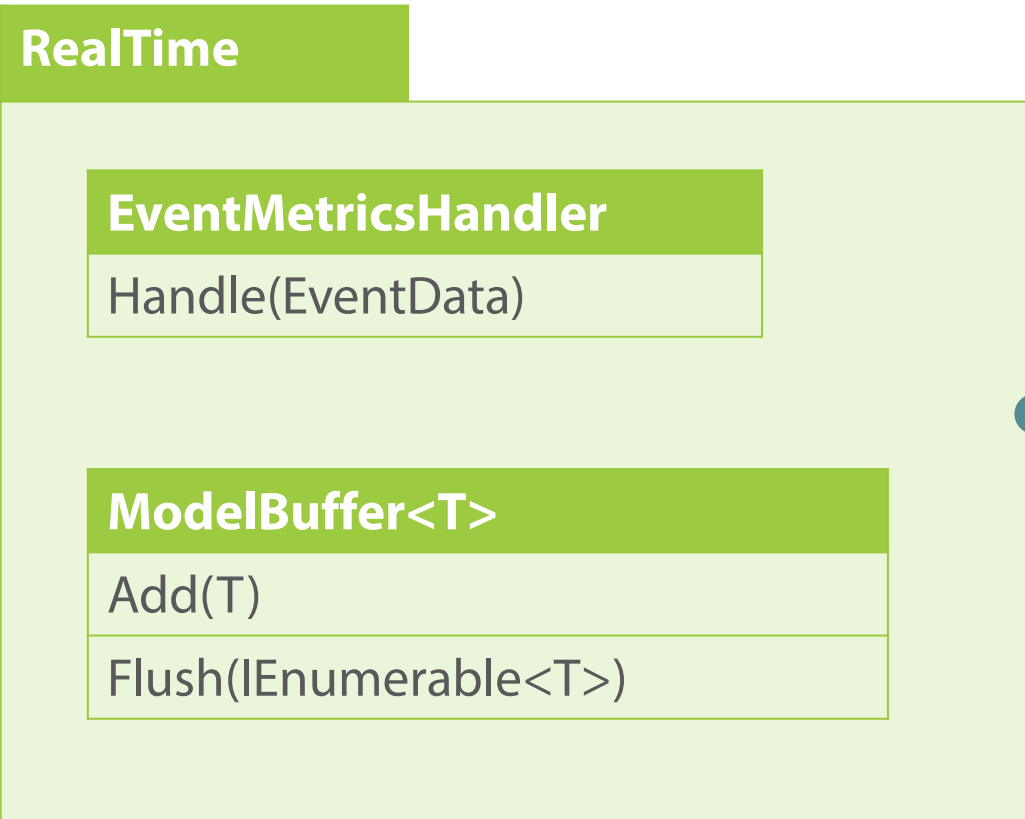
## WorkerRole


Register dependencies with Unity

```
container.RegisterType<EventReceiver>(
    new InjectionConstructor(
        new ResolvedParameter<IEventProcessorFactory>(),
        eventHubName, consumerGroupName,
        eventHubConnectionString, checkpointConnectionString));
_receiver = container.Resolve<EventReceiver>();
```

## EventReceiver

From Unity with dependency chain resolved

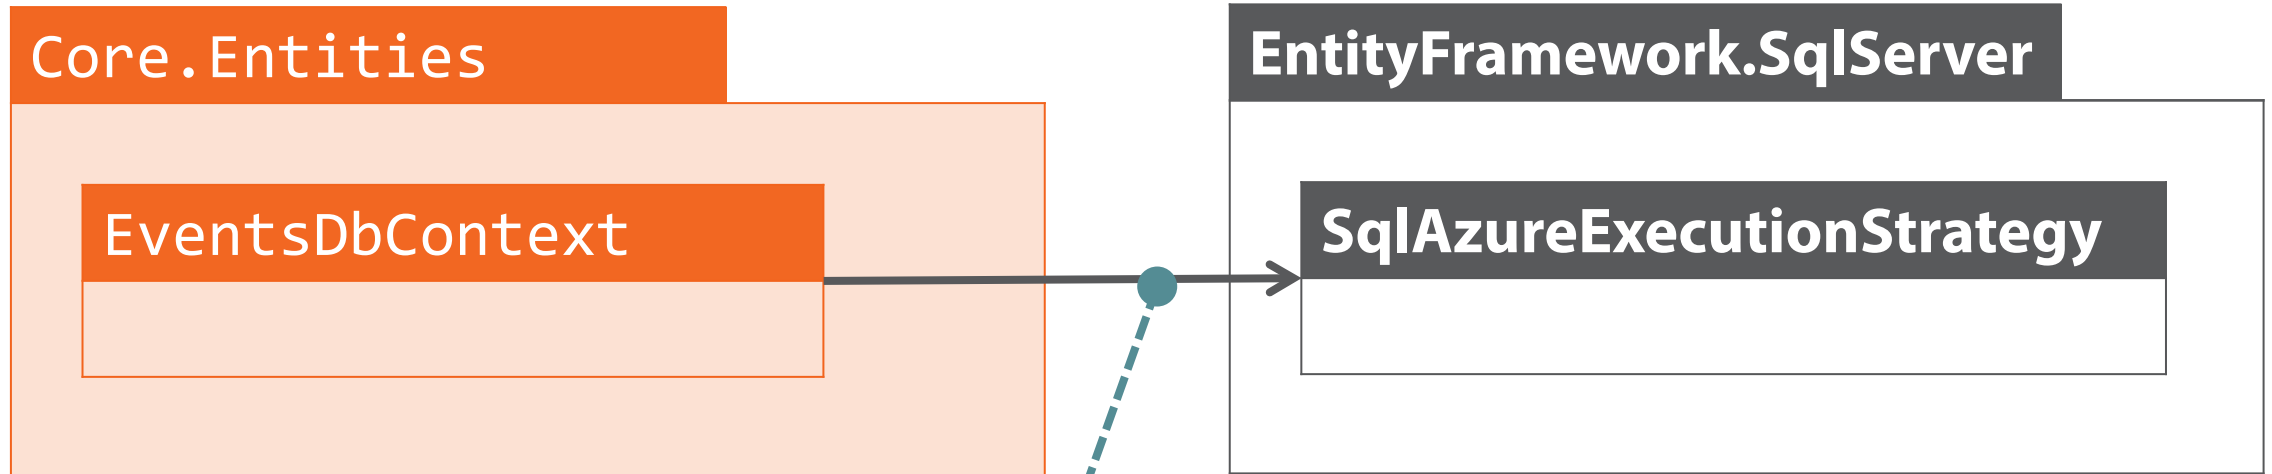






1Gbps





**Max 5 retries**  
**Max 5s wait**

**EntityFramework.SqlServer**


**SqlAzureExecutionStrategy**





**EntityFramework.SqlServer**

**SqlAzureExecutionStrategy**



# Demo: Running in Azure

Using SqlAzureExecutionStrategy

Script SQL Azure Deployment

Run RealTime worker in Azure



```
public EventsDbConfiguration()  
{  
    SetExecutionStrategy("System.Data.SqlClient", () =>  
new SqlAzureExecutionStrategy(10, TimeSpan.FromSeconds(5)));  
}
```

## EventsDbConfiguration

Inherits DbConfiguration sets SqlAzureExecutionStrategy

```
<entityFramework codeConfigurationType=  
"Telemetry.Entities.EventsDbConfiguration,Telemetry.Entities"  
>
```


## App.config



Specify codeConfigurationType for entityFramewrok


```
New-AzureSqlDatabaseServer -AdministratorLogin $AdminUser ...  
  
New-AzureSqlDatabase -DatabaseName 'eventsdb-prd' ...  
  
& "... \SqlPackage.exe" /SourceFile:"$DacPacPath"  
    /TargetConnectionString:"$connectionString" /Action:Publish
```

## Publish to SQL Azure


PowerShell Azure cmdlets and SqlPackage tool





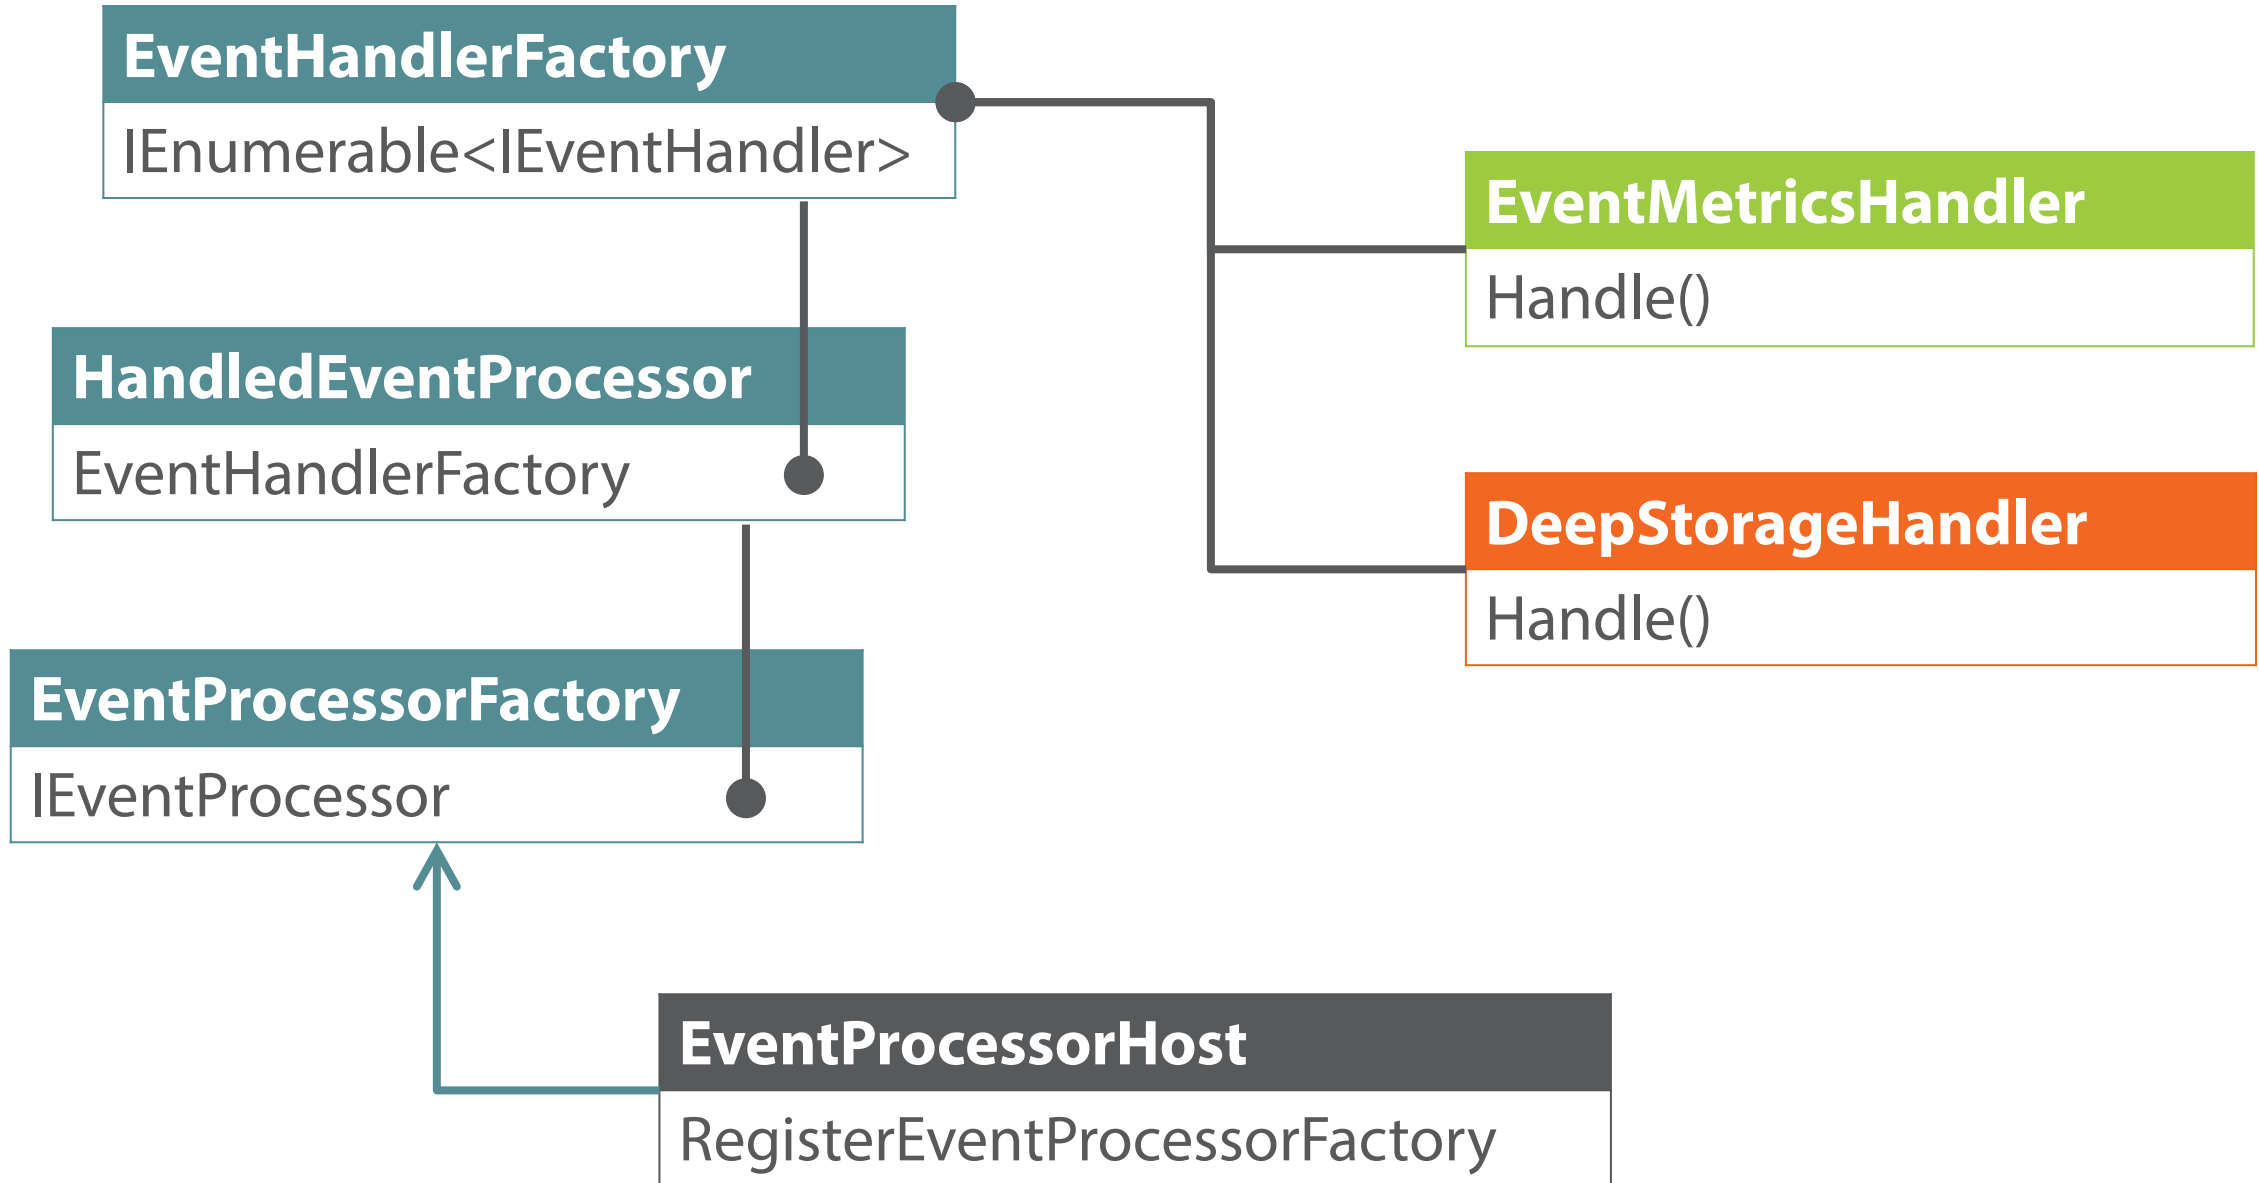


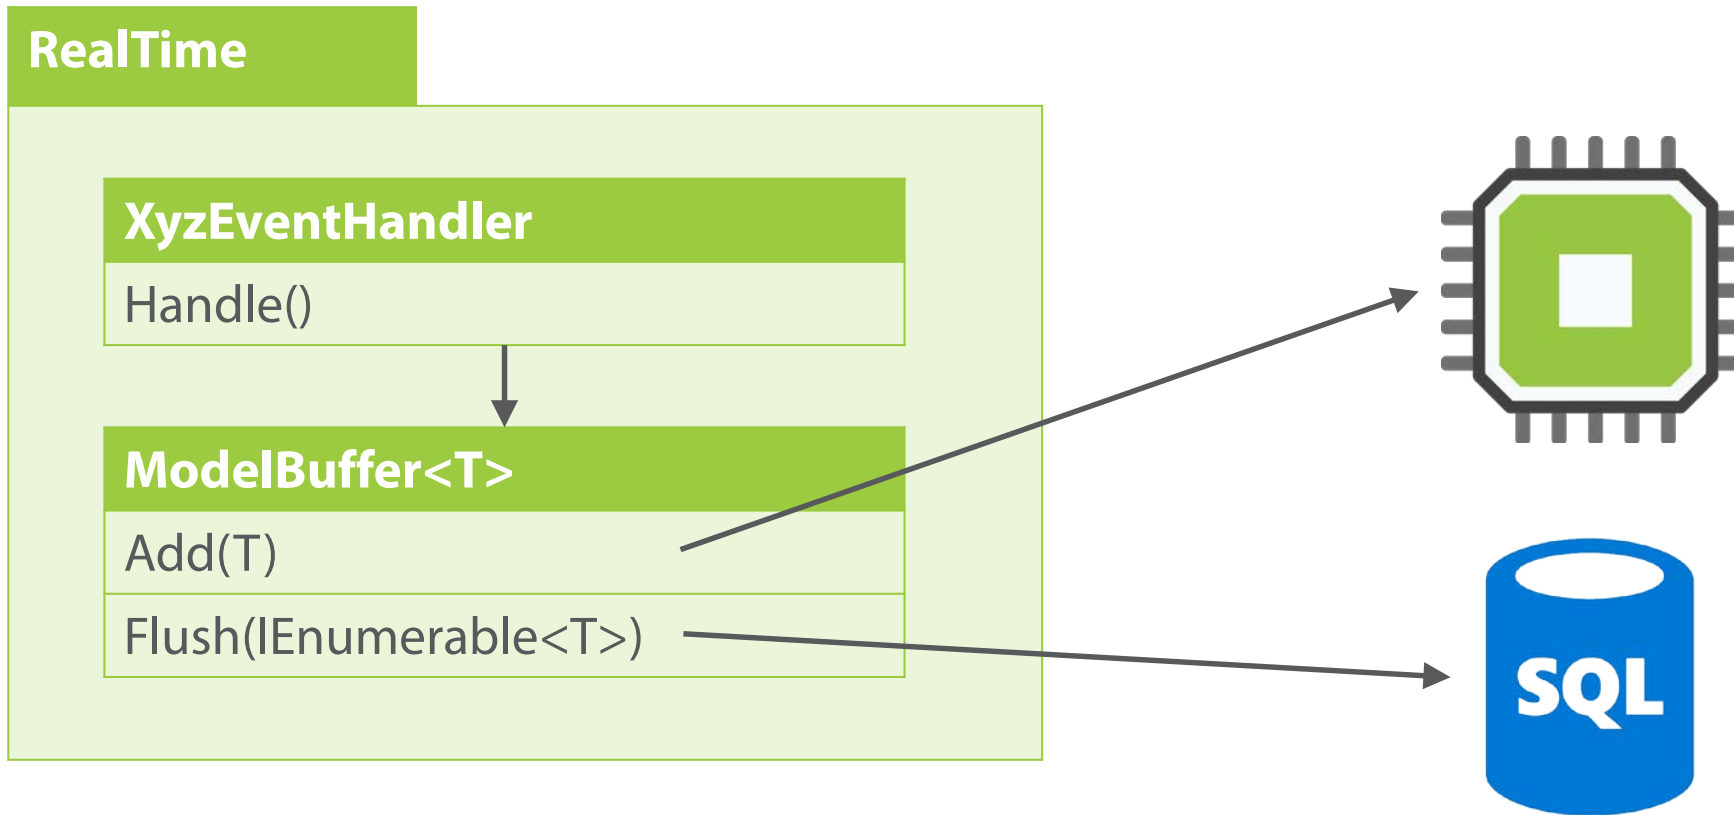




## Core.EventProcessor

- EventReceiver
- HandledEventProcessor
- IEventHandler






## EventMetricsHandler


Handle()



Event Breakdown



device.activated	device.media.newmedia	system.app.focus.lost	device.orientation.changed
system.app.updated	device.metrics.sampled	device.logs.message	system.user.created
device.usage.started	device.usage.stopped	device.wifi.connected	system.app.focus.gained



**ClientActivityHandler**

Handle()



DeviceId	TotalUsage	UsageThisWeek	UsageThisMonth	UsageThisYear
a6b56d5b64	6,791	520	1,224	6,791
cb4375f2aa	11,203	1,604	4,376	9,543
bace33211e	102	102	102	102
dff1344fee	11	1	10	11

